Amazon EMR

What is Amazon EMR

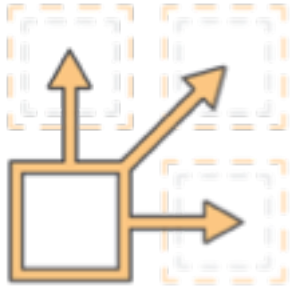

Easy to use
Launch a cluster in minutes

Low cost
Pay per-second

Open-source variety
Latest versions of software

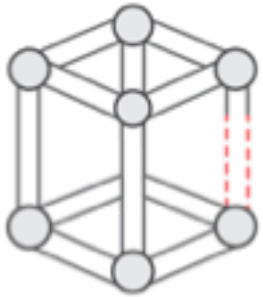

Managed
Spend less time monitoring

Secure
Easy to enable options

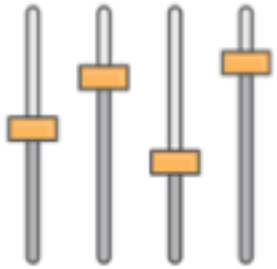

Flexible
Full customization and control

aws

Enterprise-grade Hadoop & Spark

Deploy latest releases in Hadoop and Spark ecosystems

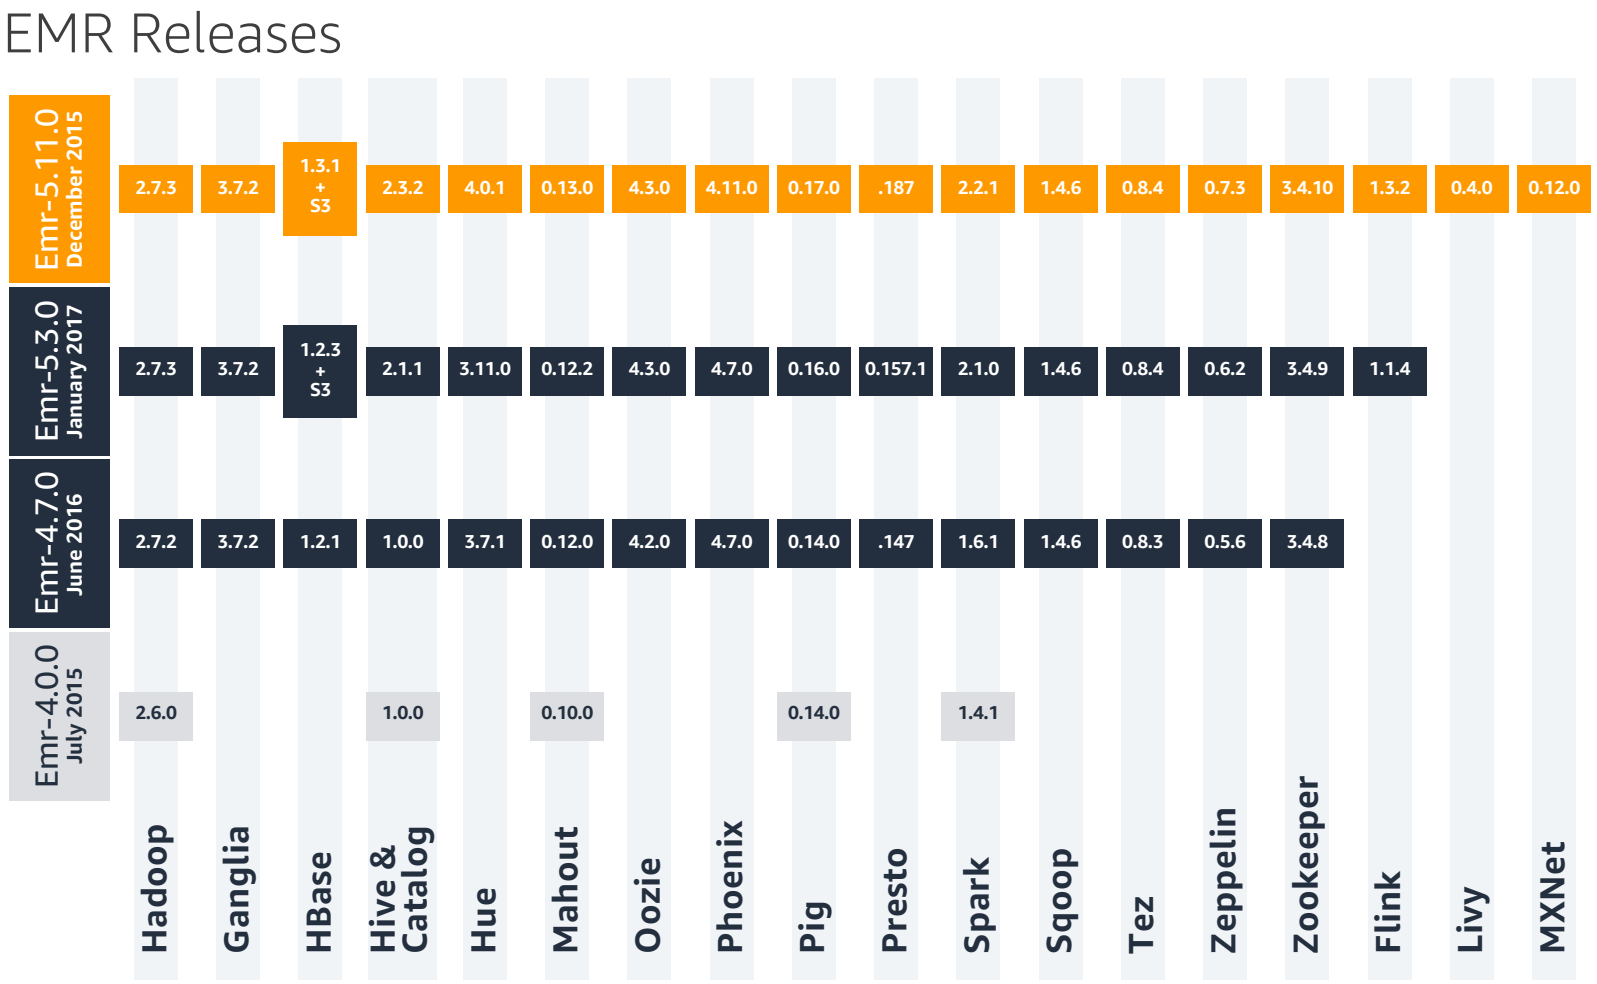

- 19 open-source projects: Apache Hadoop, Spark, HBase, Presto, and more
- Updated with the latest open source frameworks within 30 days of release

aws

Open-source applications

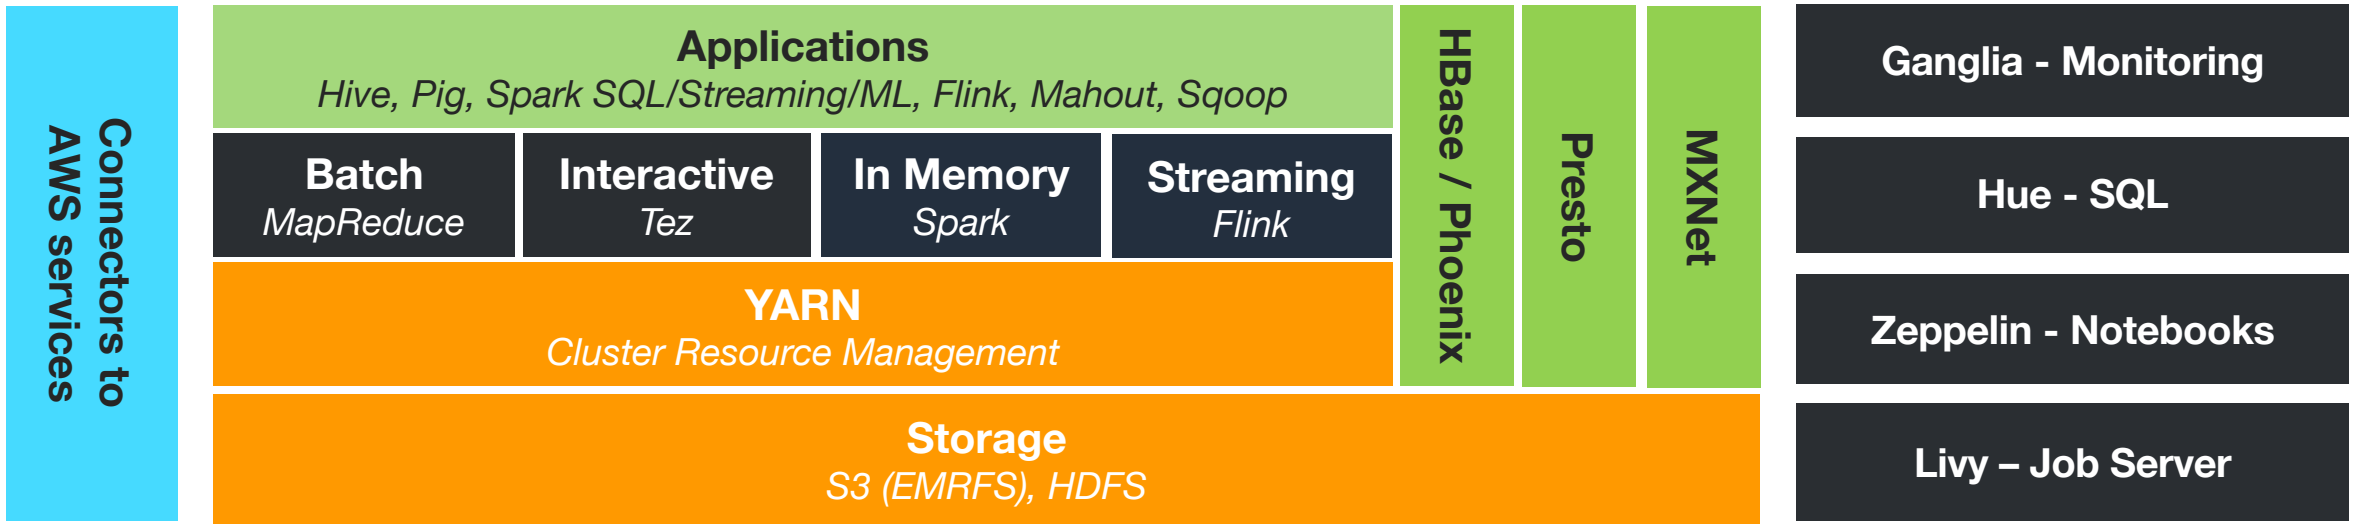

Amazon EMR service

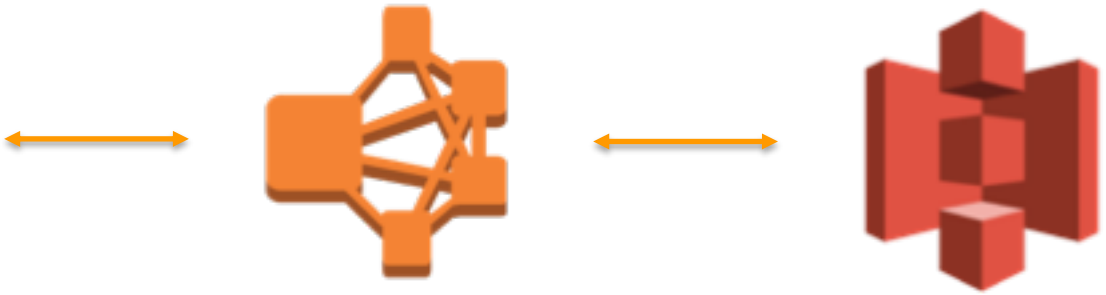

Enterprise-grade Hadoop & Spark

Scale to any size

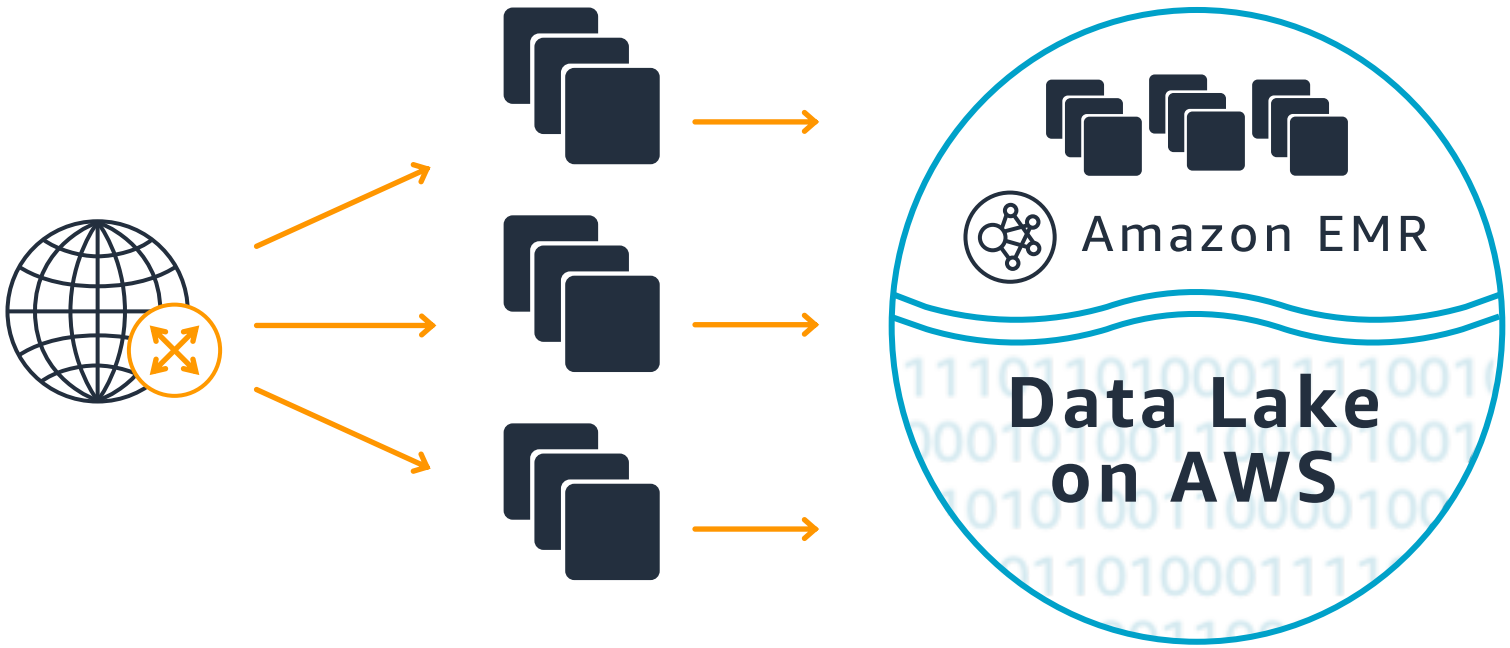

- Scale compute (EMR) & storage (S3) independently
- Store, and process any amount of data—PB to EBs
- Provision one, hundreds, or thousands of nodes
- Auto-scaling

Enterprise-grade Hadoop & Spark

Highly available and durable

- S3 is designed to deliver 99.999999999% durability
- EMR monitors your cluster—replacing poorly performing & failed nodes, and restarting services
- Monitor your clusters using Amazon CloudWatch
- Built-in console to view job history & browse logs
- EMR has on-cluster HDFS for data persistence

aws

Enterprise-grade Hadoop & Spark

Highly secure

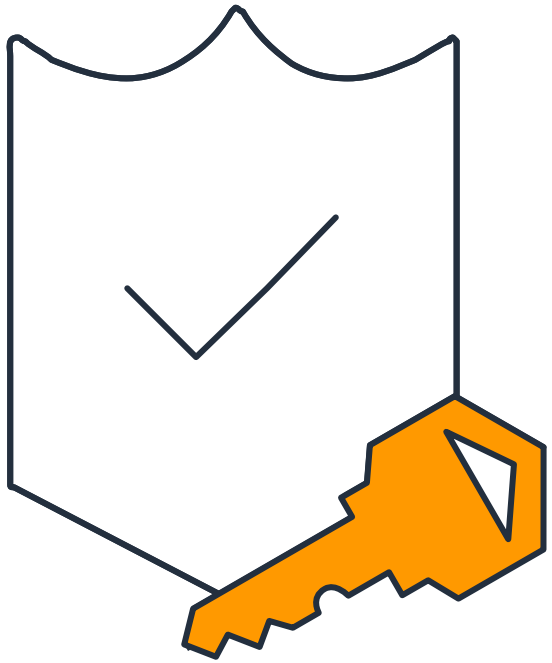

- Encryption of data at rest and in-transit
- ML-powered security with Amazon Macie
- Network isolation using Amazon VPC
- Access and permissions control with IAM policies
- Log, and audit activity with AWS CloudTrail
- Microsoft AD integration with Kerberos support

aws

Options to submit jobs

Submit a Spark application

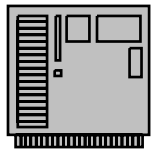

Amazon EMR Step API

Use AWS Lambda to submit applications to EMR Step API or directly to Spark on your cluster

AWS Lambda

AWS Data Pipeline

Create a pipeline to schedule job submission or create complex workflows

Airflow, Luigi, or other schedulers on EC2

Amazon EMR

Use Oozie on your cluster to build DAGs of jobs

Networking: VPC private subnets

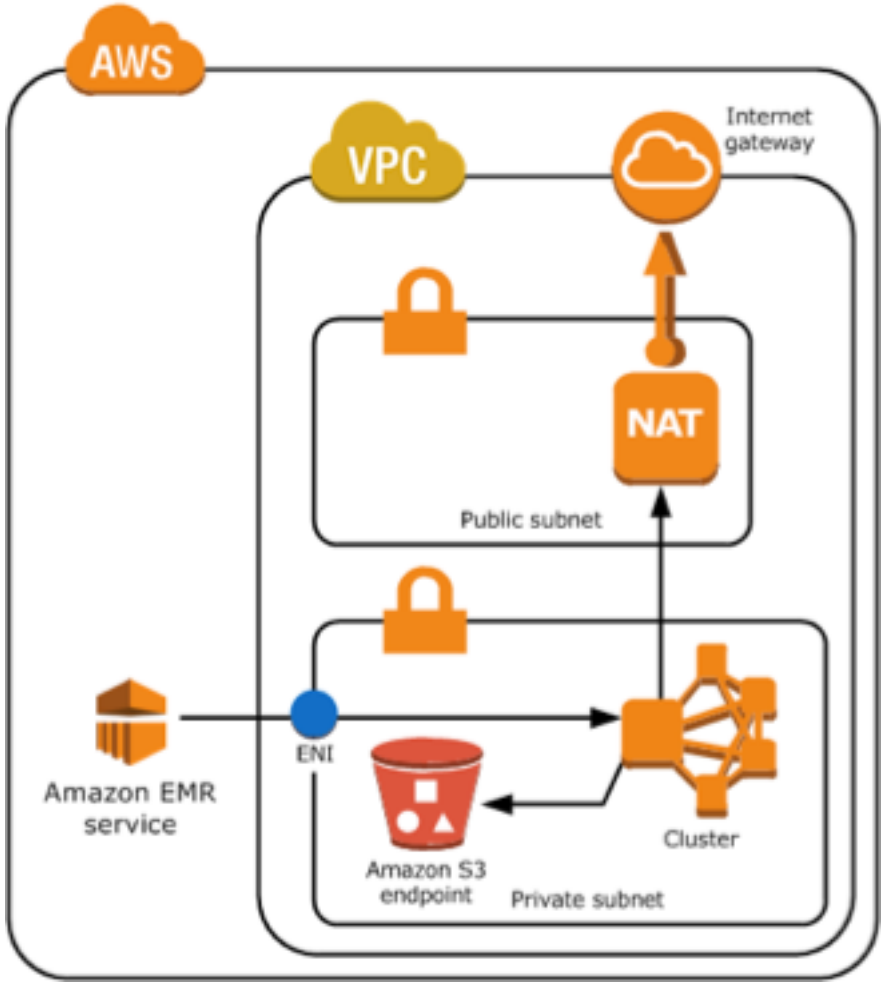

- Use Amazon S3 Endpoints for connectivity to S3
- Use Managed NAT for connectivity to other services or the Internet
- Control the traffic using Security Groups
  - ElasticMapReduce-Master-Private
  - ElasticMapReduce-Slave-Private
  - ElasticMapReduce-ServiceAccess

Access Control: IAM Users and Roles

- IAM Policies for access to Amazon EMR service (IAM users or federated users)
  - **AmazonElasticMapReduceFullAccess**
  - **AmazonElasticMapReduceReadOnlyAccess**

- IAM Policies for Amazon EMR cluster
  - Service role (**AmazonElasticMapReduceRole)** - Allowable actions for Amazon EMR service, like creating EC2 instances.
  - Instance profile (**AmazonElasticMapReduceforEC2Role)** - Applications that run on Amazon EMR, like access to Amazon S3 for EMRFS on your cluster.

Easy to Configure End-to-End Security

- Encryption for data at rest
  - Process encrypted data from Amazon S3 with support for all Amazon S3 encryption features
  - Configurable Local disk and HDFS encryption
- Encryption for data in transit
  - Run EMR clusters in VPC private subnets
  - Encrypted inter-node communication for Hadoop, MapReduce, Spark
  - Data transfer to other services over SSL
- Integrated with AWS IAM
  - Support for IAM roles, Bucket policies & ACLs, Tag-based permissions
- Authentication and authorization with native Hadoop ecosystem feature set
- Compliance and Auditing
  - SOC 1/2/3, PCI-DSS, FedRAMP, HIPAA-eligible
  - All API calls logged in CloudTrail
  - Object access logging for S3 data

aws

Encryption – use security configurations

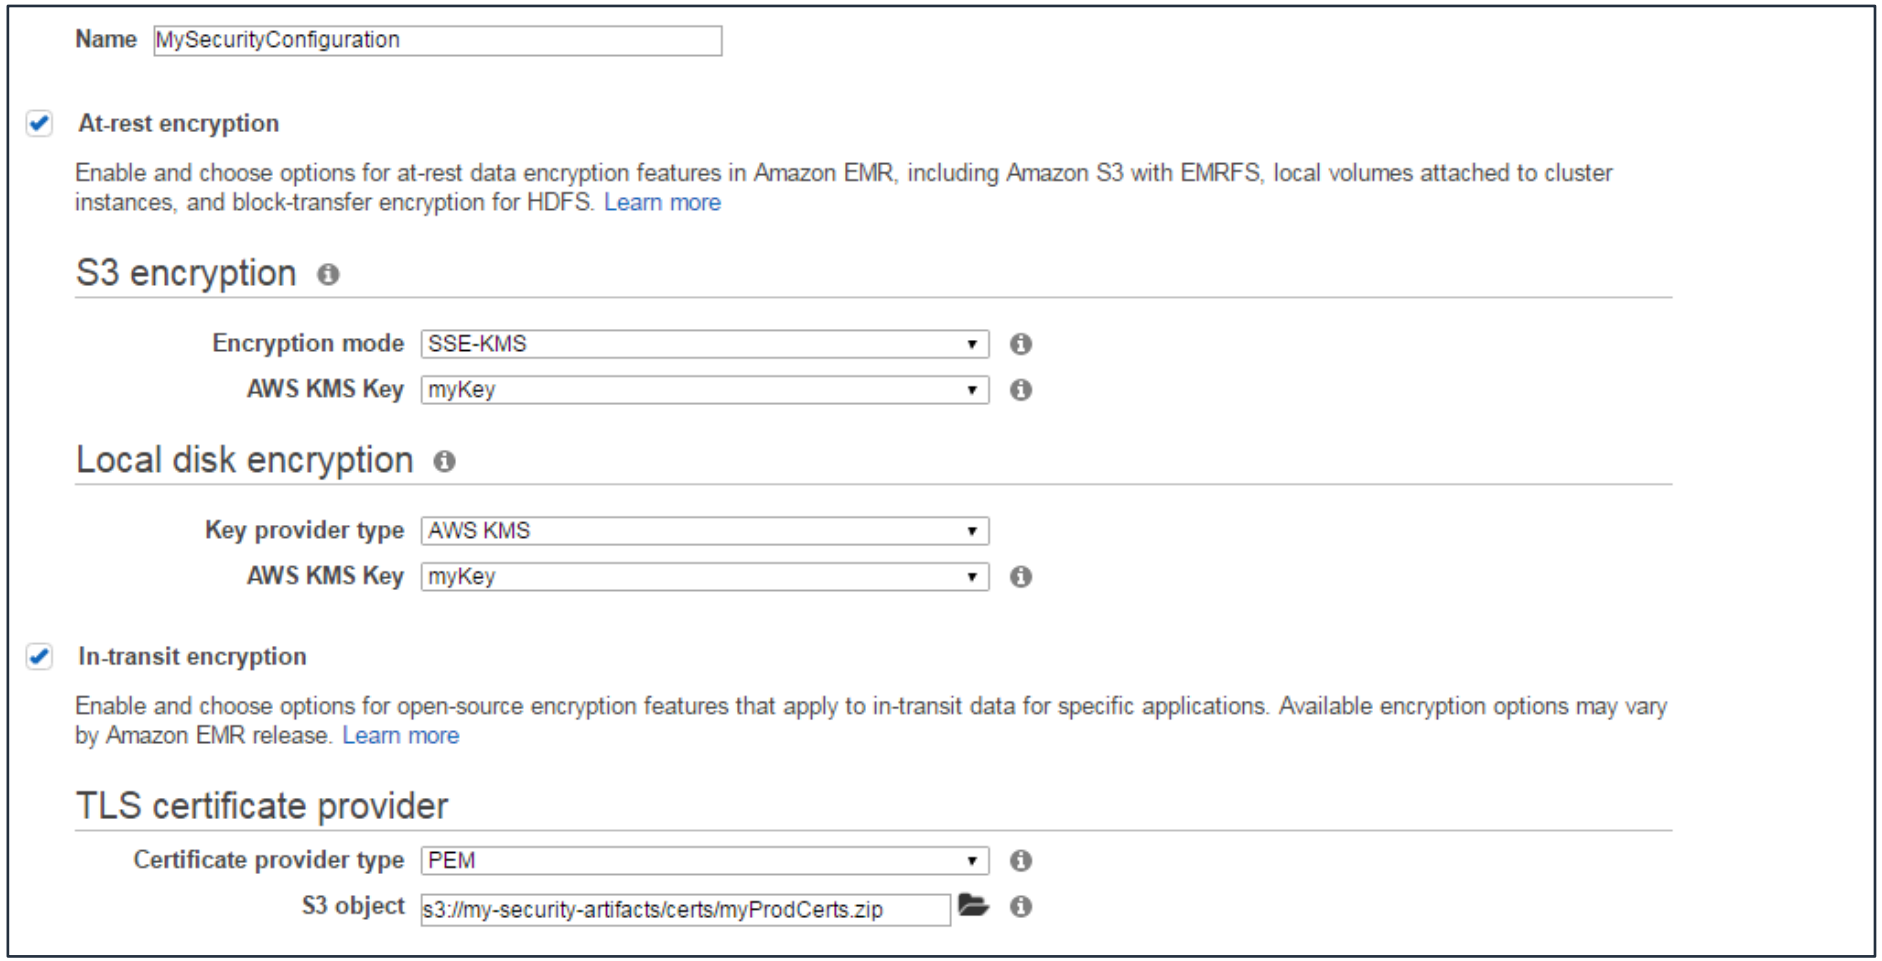

Data at Rest: S3 client-side encryption

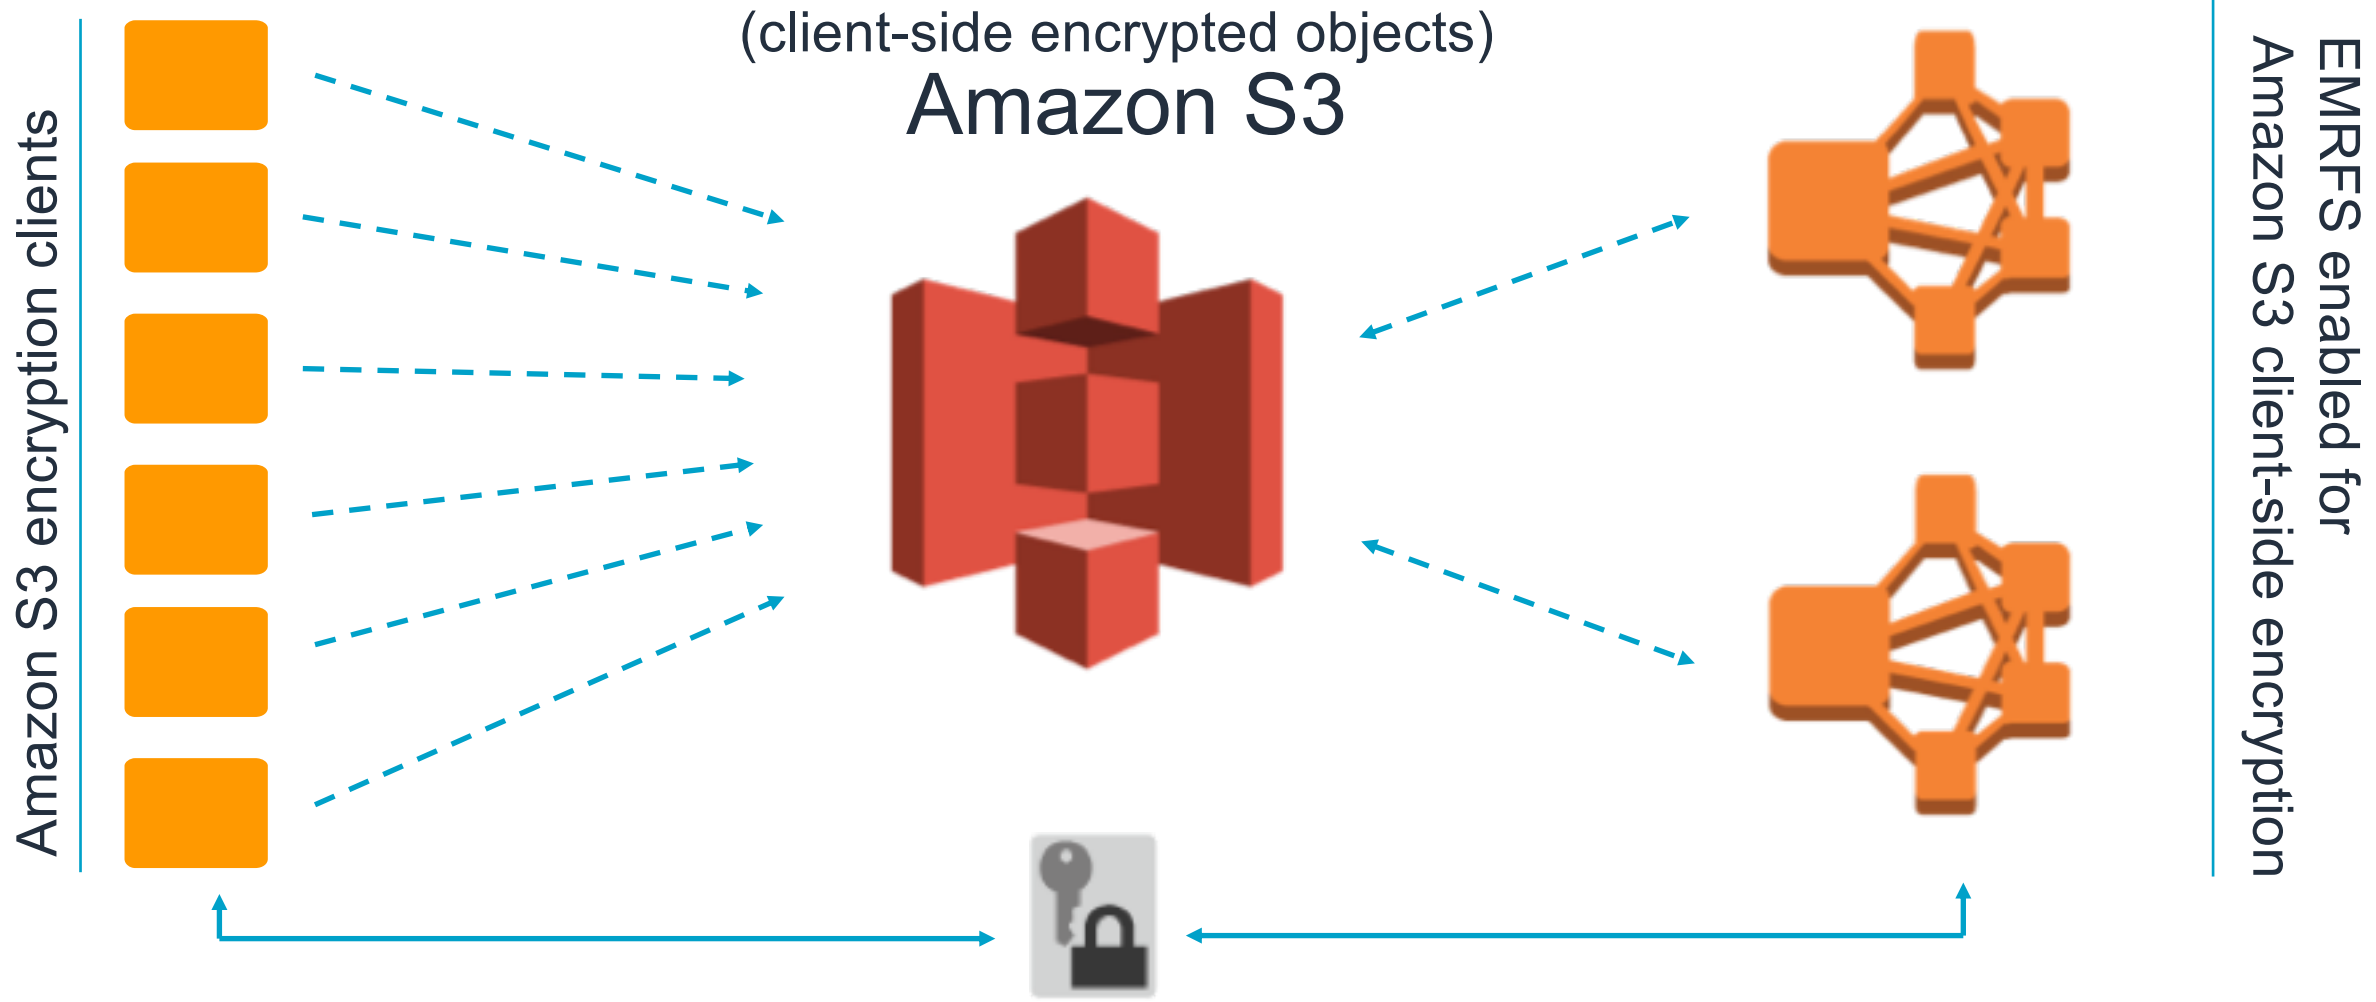

Security - Encryption

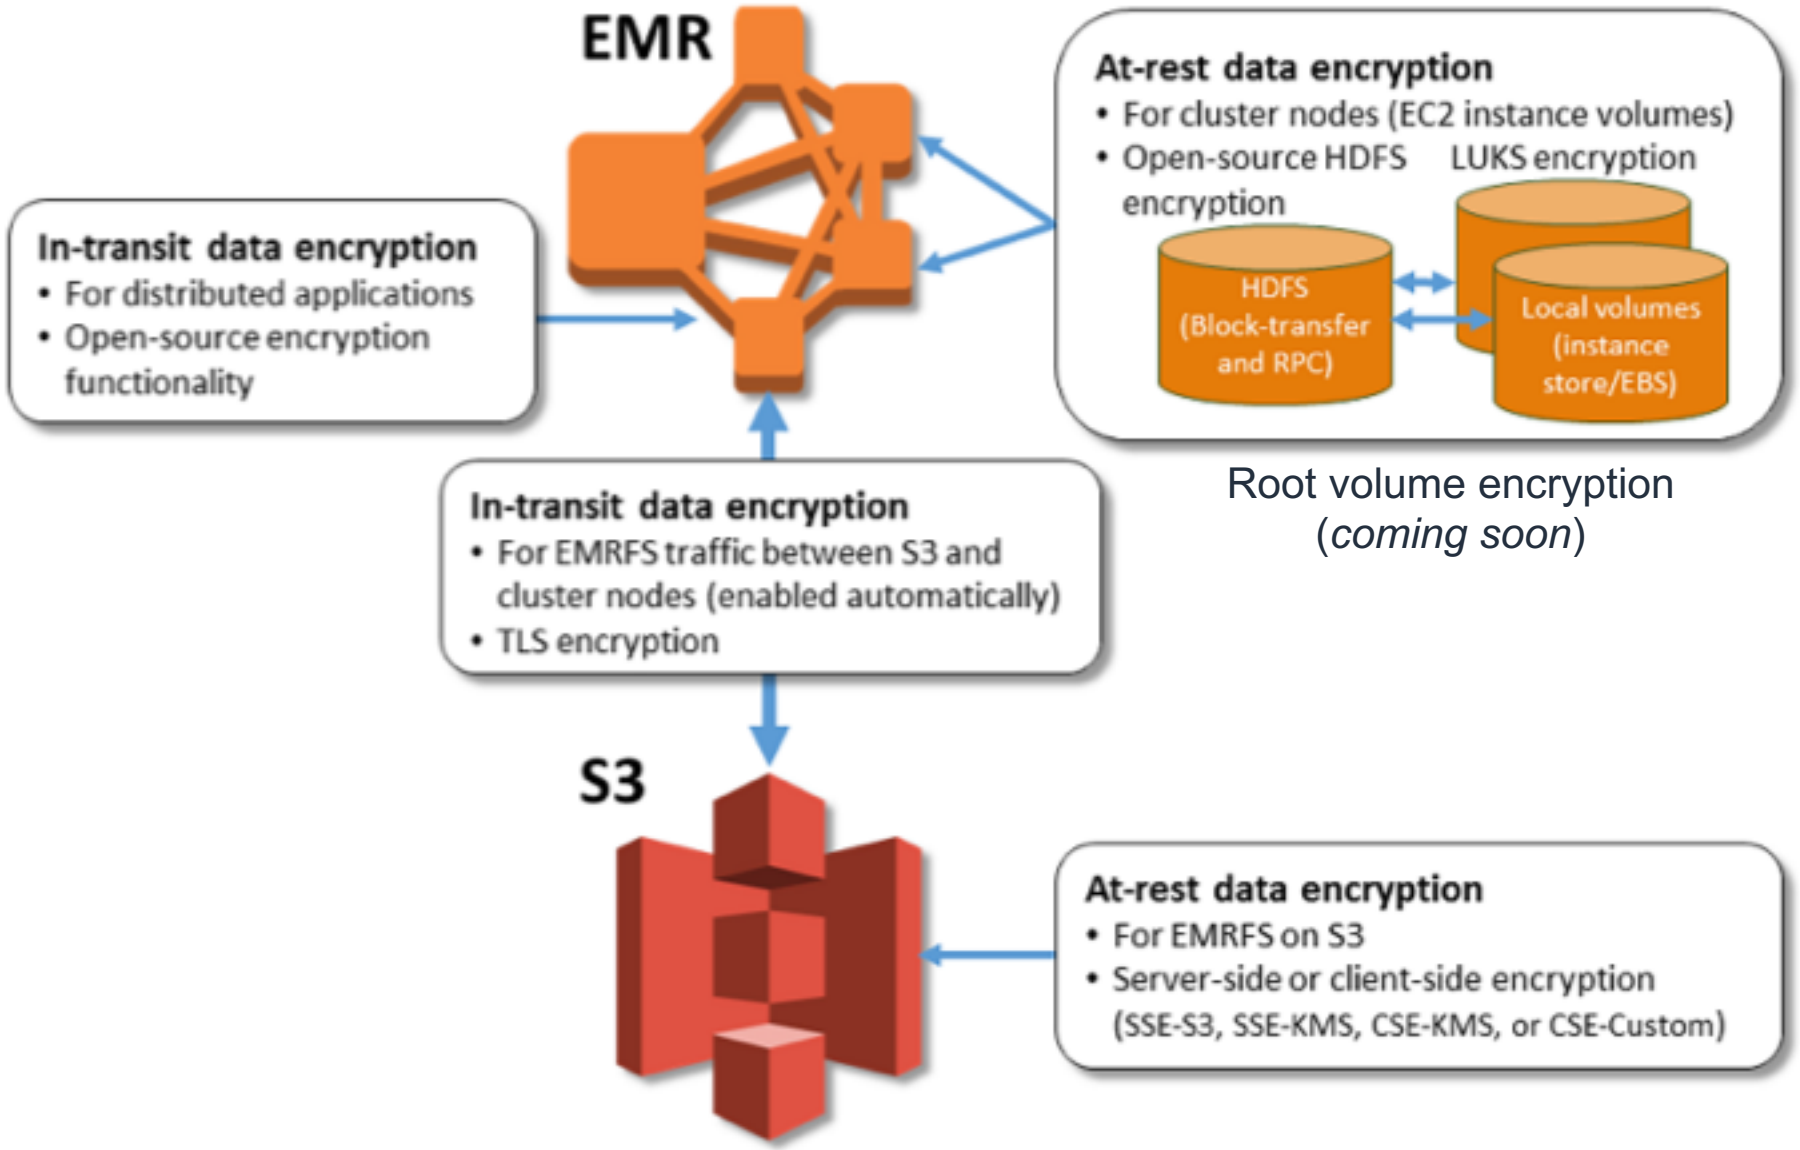

Root volume encryption
(*coming soon*)

Supported
- Spark
- Tez
- MapReduce

Security – Authentication and Authorization

EMR role
EC2 role
SSH key

IAM user: `MyUser`

Tag: `user = MyUser`

```
{
    "Sid": "Stmt1479329681000",
    "Effect": "Allow",
    "Action": [
        "elasticmapreduce:AddTags",
        "elasticmapreduce:RunJobFlow"
    ],
    "Condition": {
        "StringEquals": {
            "elasticmapreduce:RequestTag/user": "MyUser"
        }
    },
    "Resource": [
        "*"
    ]
}
```

```
        "Effect": "Allow",
        "Action": [
            "elasticmapreduce:AddJobFlowSteps",
            "elasticmapreduce:DescribeCluster",
            "elasticmapreduce:DescribeStep",
            "elasticmapreduce:ListSteps",
            "elasticmapreduce:TerminateJobFlows"
        ],
        "Condition": {
            "StringEquals": {
                "elasticmapreduce:ResourceTag/user": "MyUser"
            }
        },
        "Resource": [
            "*"
        ]
    }
```

Security - Authentication and Authorization

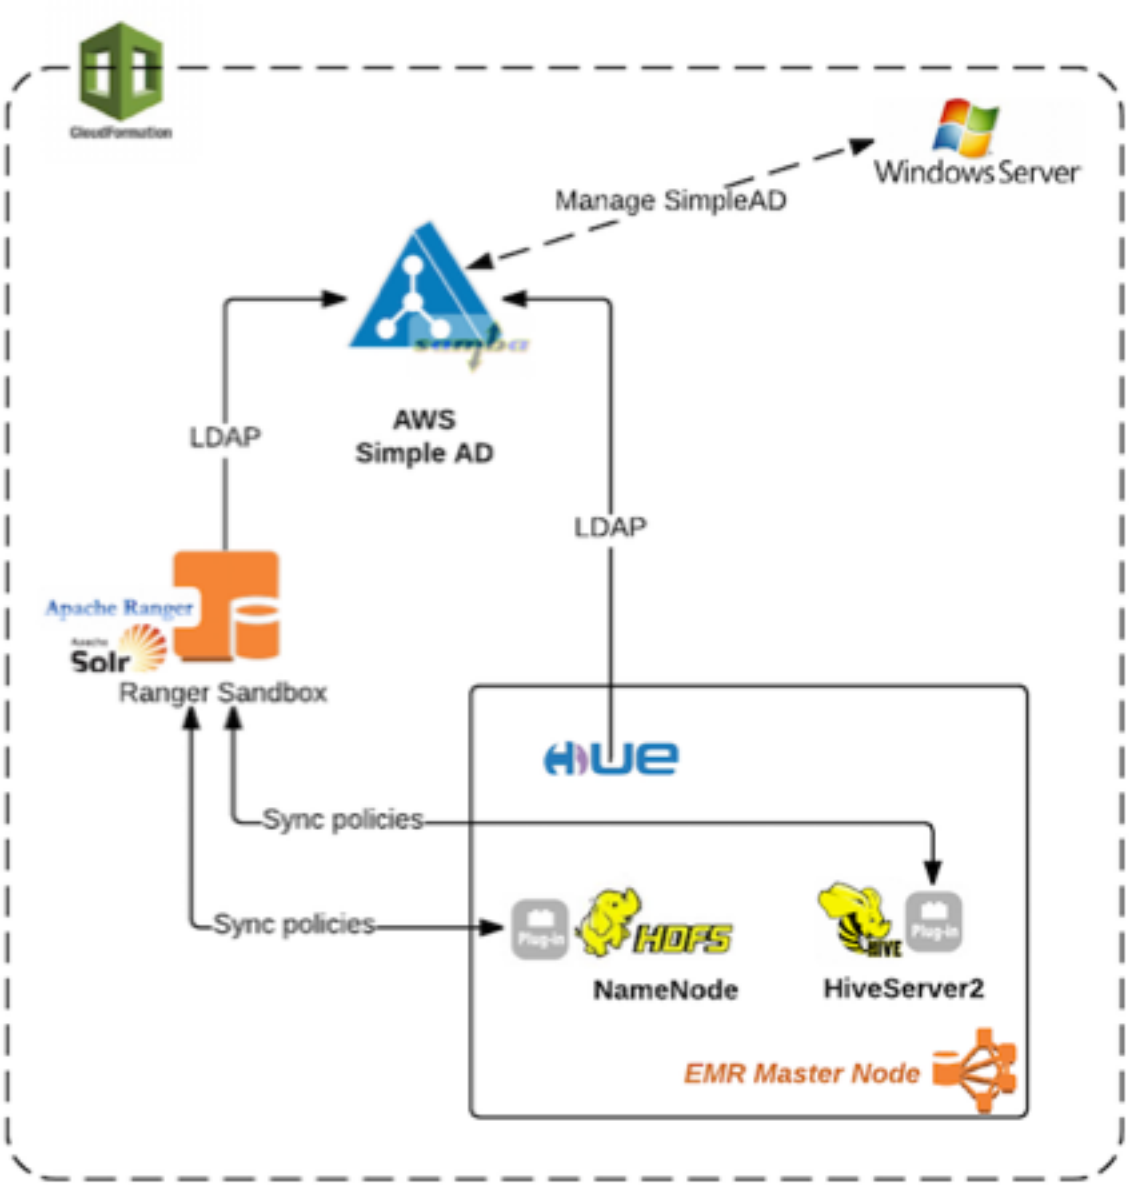

Apache Ranger

- Plug-ins for Hive, HBase, YARN, and HDFS
- Row-level authorization for Hive (with data-masking)
- Full auditing capabilities with embedded search
- Run Ranger on an edge node – visit the AWS Big Data Blog

Authentication

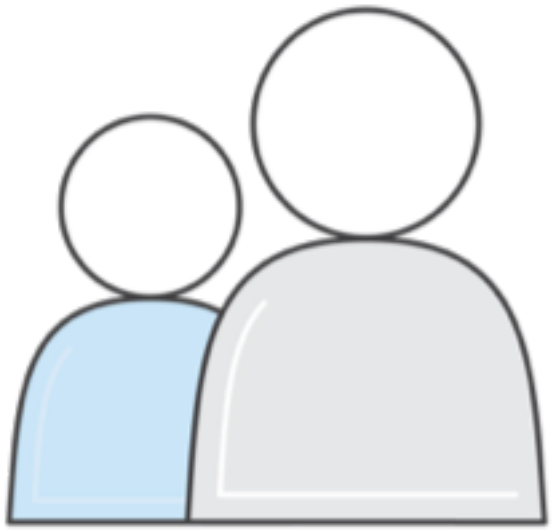

LDAP
HiveServer2
Presto Coordinator
Spark Thrift Server
Hue Server
Zeppelin Server

EC2 key pair
SSH as “hadoop”

AWS credentials
EMR Step (EMR API)

Authentication with Kerberos

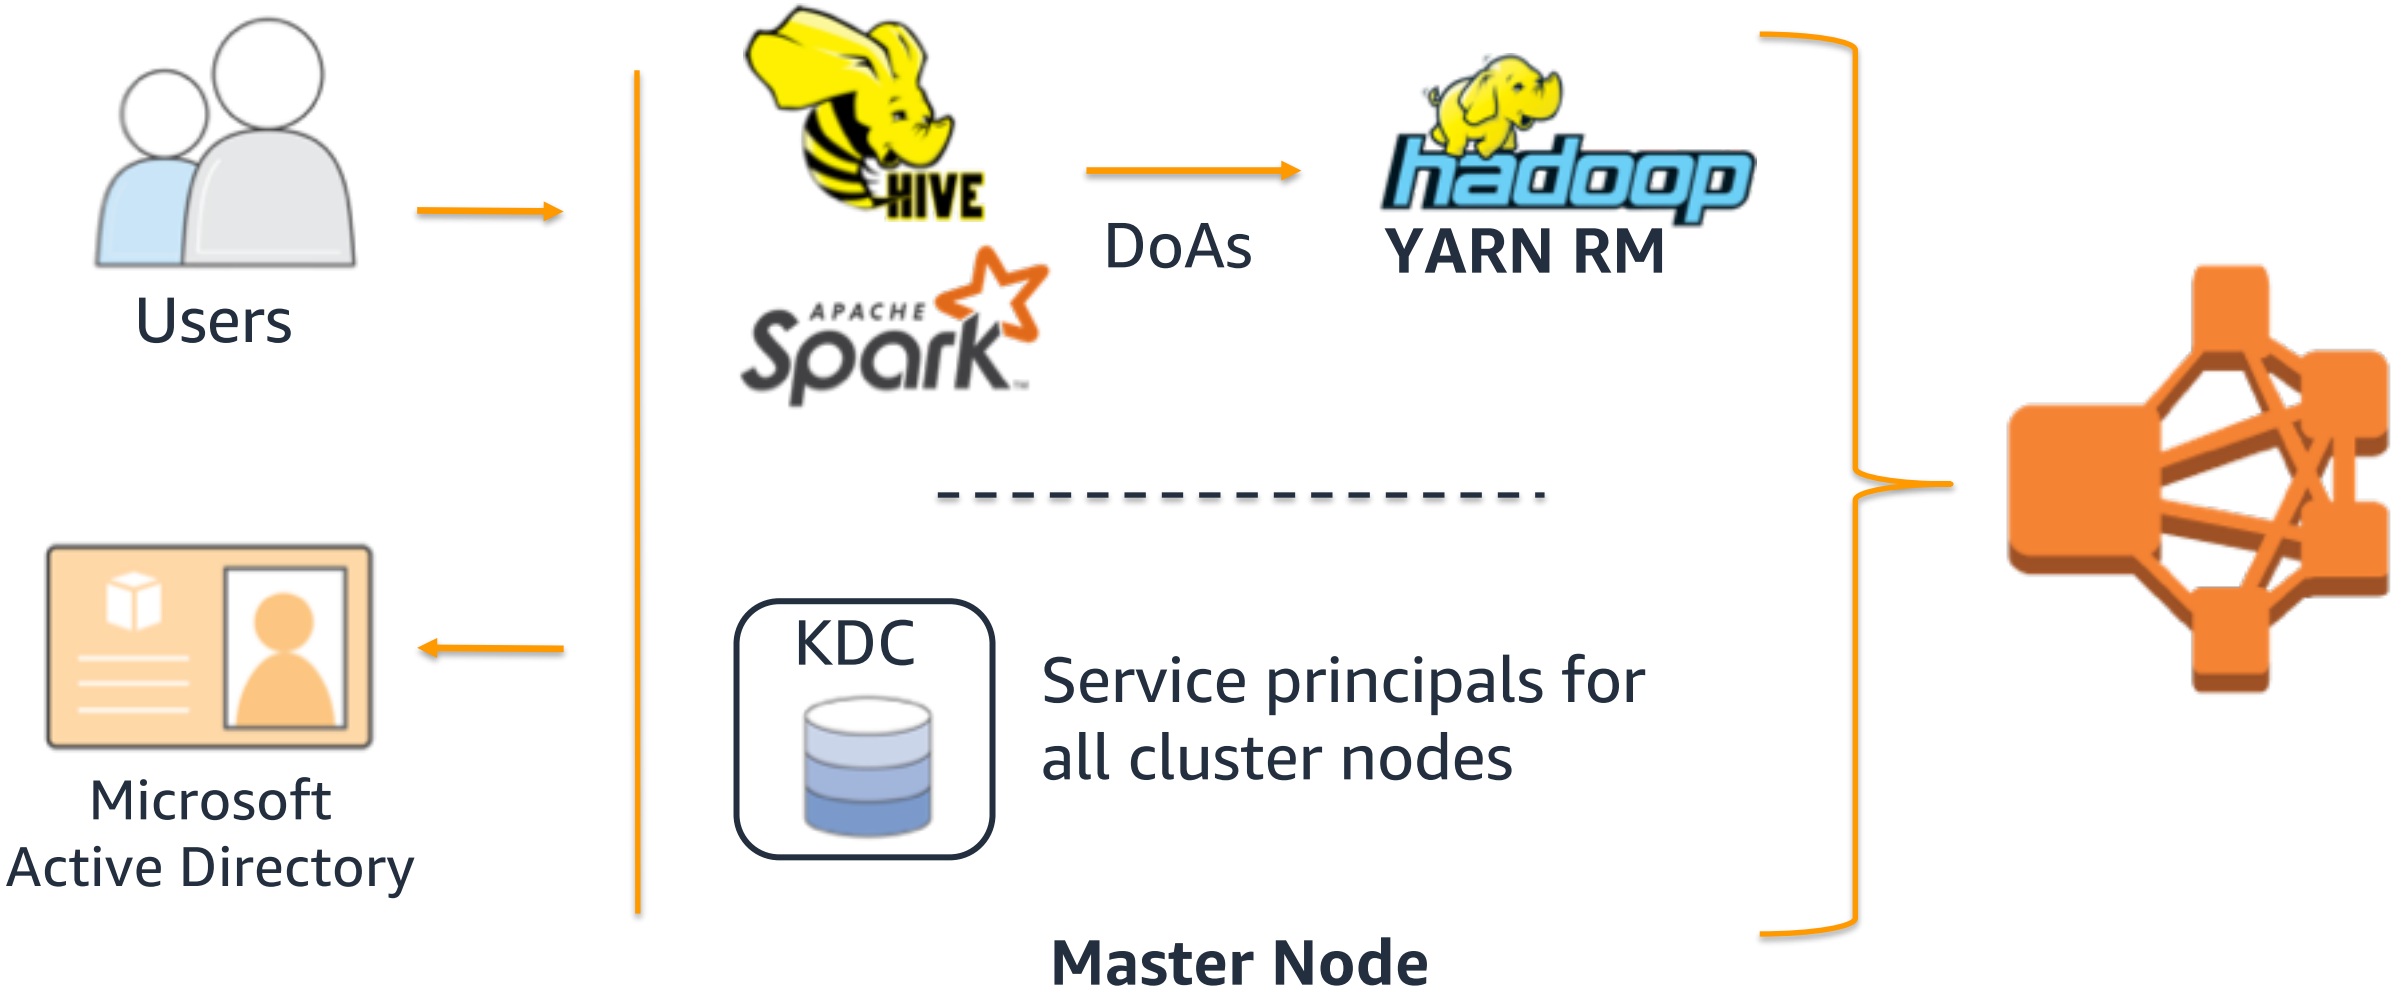

Authorization

- Storage-based
  - EMRFS/S3
  - HDFS
- HiveServer2 and Presto (SQL-based)
- HBase
- YARN queues
- Fine-grained access control by cluster tag (IAM)
- Apache Ranger on edge node (using CloudFormation)

EMRFS fine-grained authorization

Context
User: aduser
Group: analyst

IAM role: analytics_prod

Can map IAM roles to user, group, or S3 prefix

Security – Governance and Auditing

- AWS CloudTrail for EMR APIs
- S3 access logs for cluster S3 access
- YARN and application logs
- Ranger for UI for application level auditing

Monitoring

Full **visibility** of your AWS environment

- AWS CloudTrail will record access to API calls and save logs in your S3 buckets, no matter how those API calls were made

Who did **what** and **when** and from **where** (IP address)

- AWS CloudTrail support for many AWS services and growing - including **Amazon EMR**

Monitoring

- Amazon S3
  - Bucket access logs
- Amazon EMR
  - Archives various log files to Amazon S3 at 5 minute intervals.
  - Log files are available after the cluster terminates
- CloudWatch Metrics
  - Updated every five minutes and archived for two weeks
- Ganglia

Use the AWS Glue Data Catalog

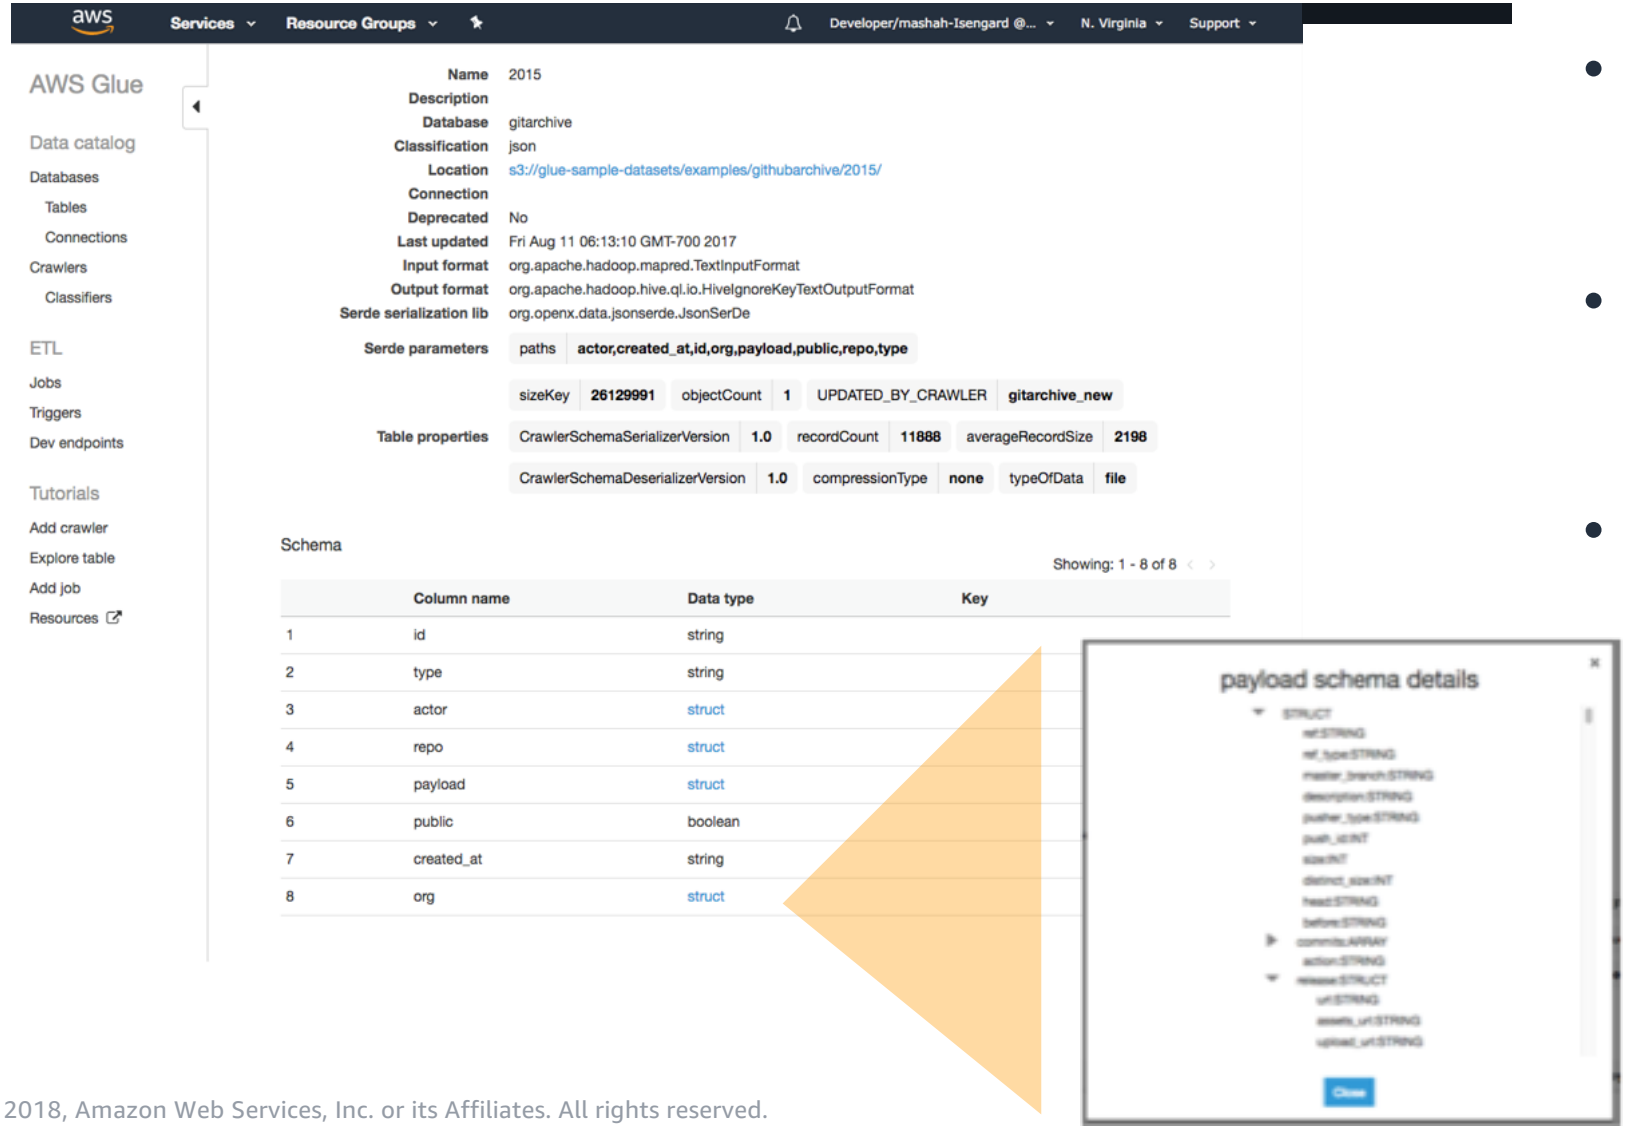

- Support for Spark, Hive and Presto
- Auto-generate schema and partitions
- Managed table updates

Real-time and batch processing

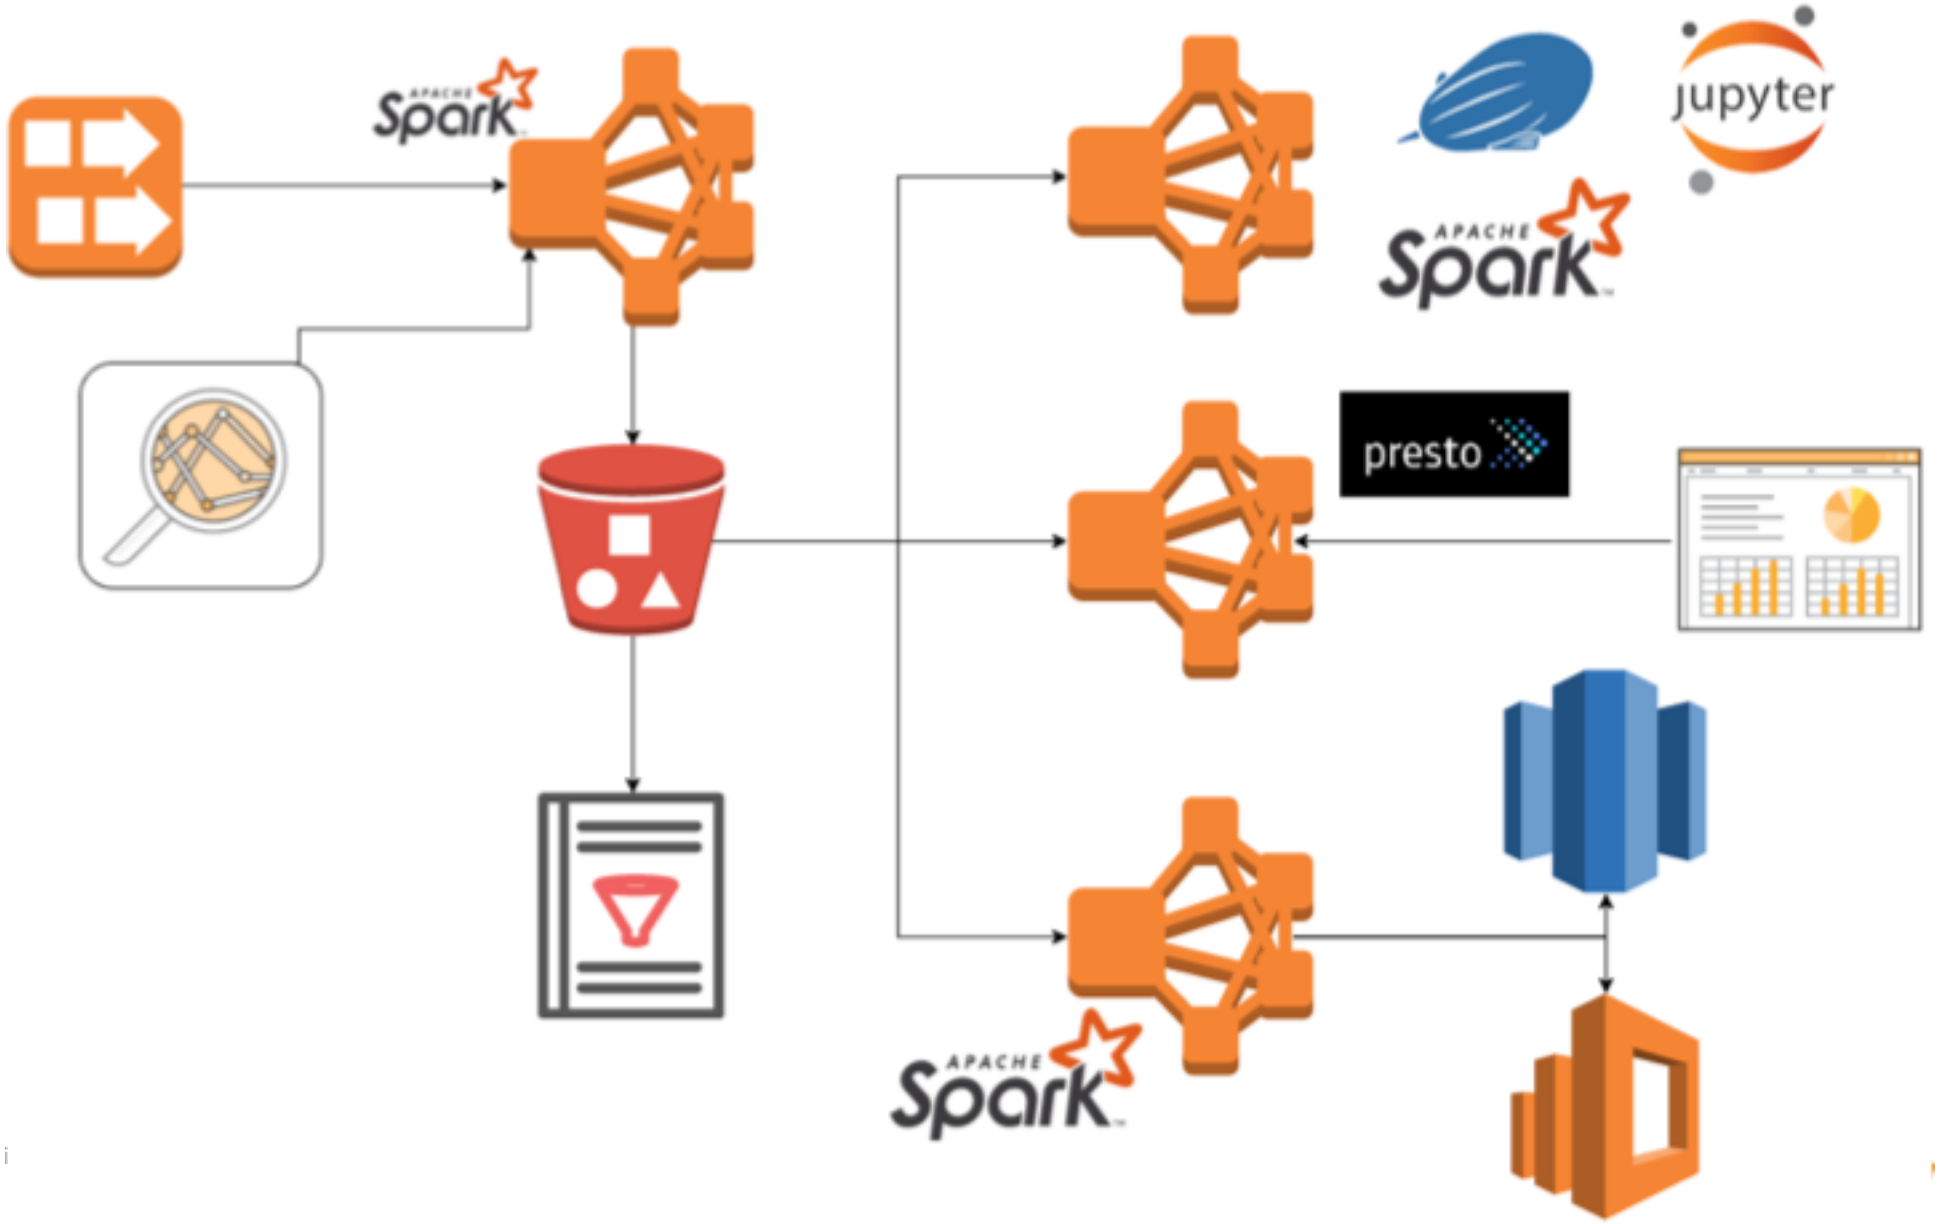

HBase for random access at massive scale

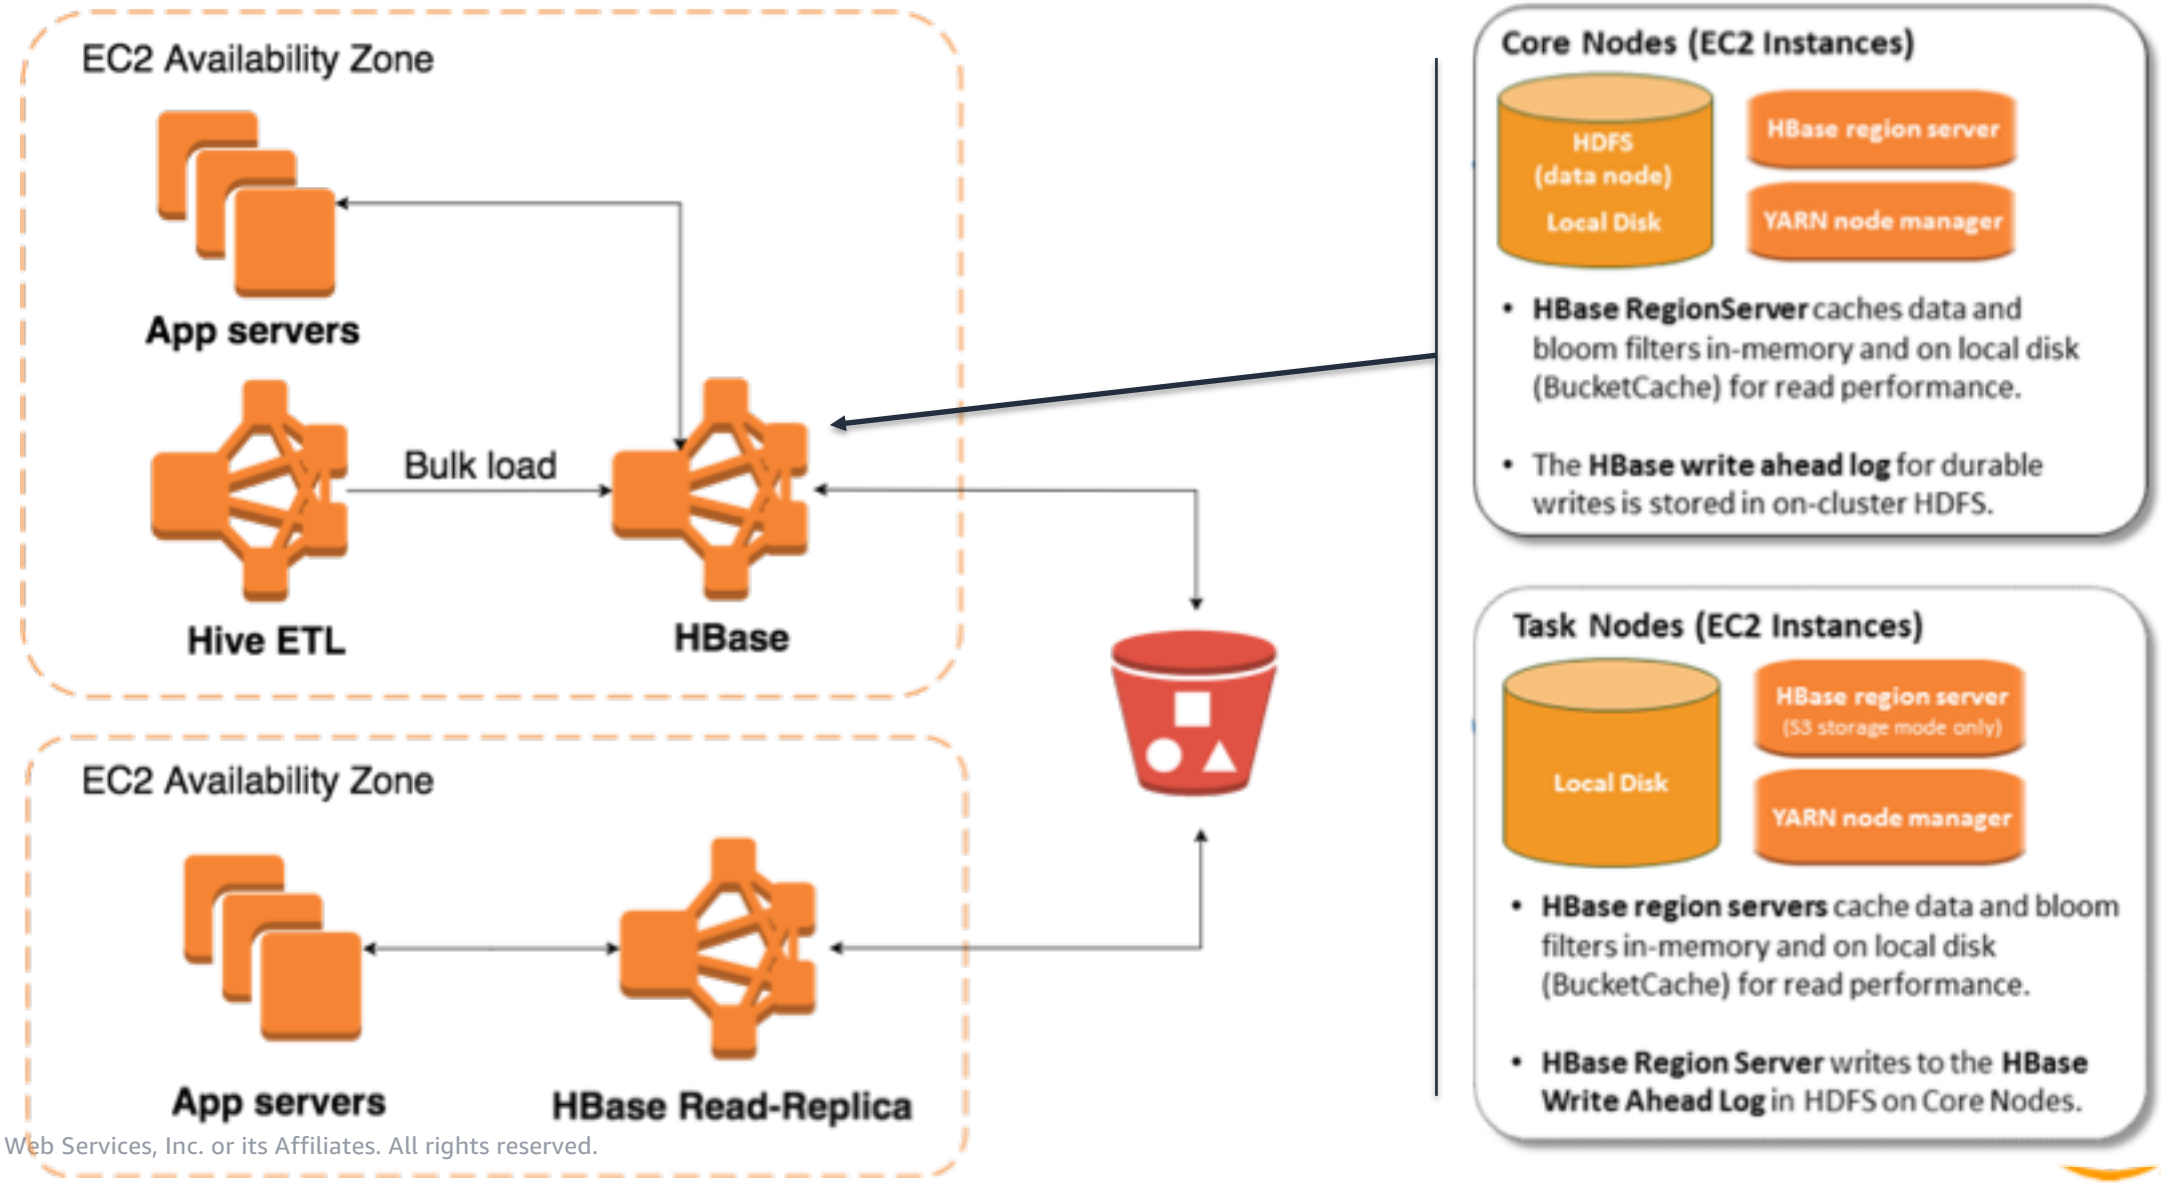

Deep learning with GPU instances

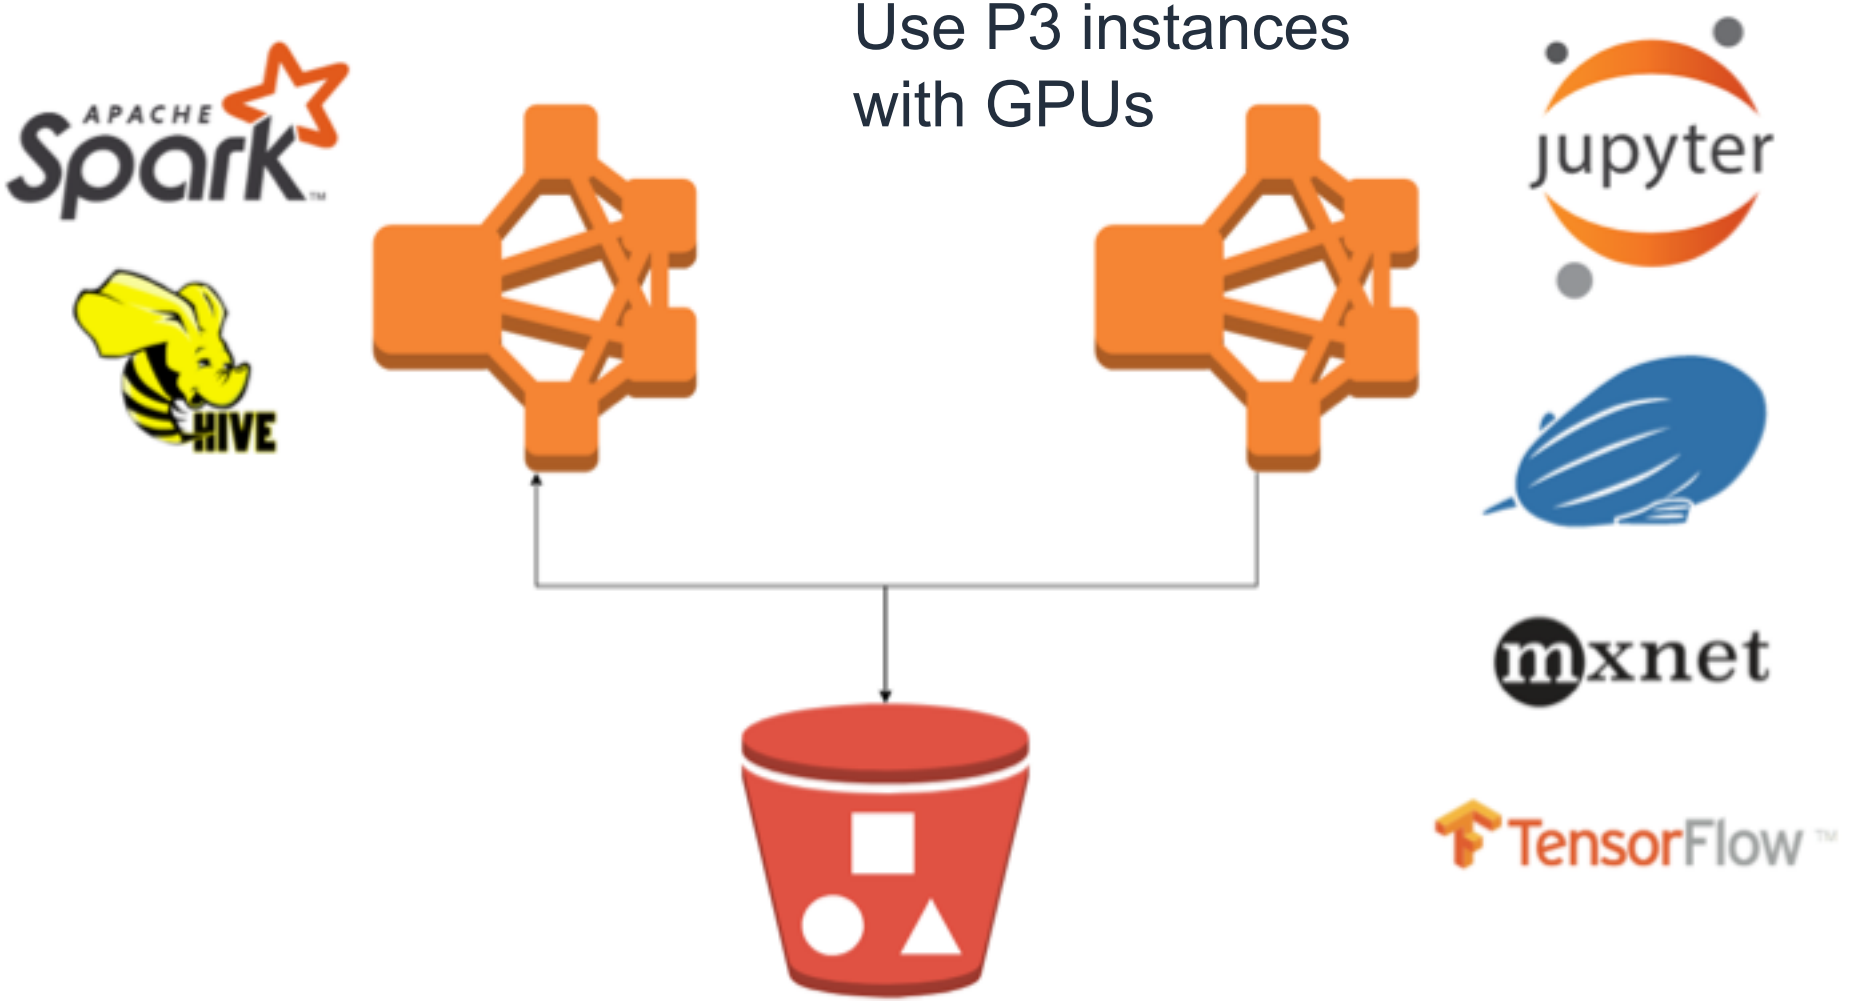

Use a custom Amazon Linux AMI

- Launch clusters with your Amazon Linux AMI
- Preinstall custom software for faster start times
- Encrypt the root volume with an AWS KMS key
- Adjust root volume size for custom applications

Application History – EMR console

Jobs > Job 0 > Stage 1 (attempt 0)

Total time across all tasks: 19.4 h
Locality level summary: Process local: 500
Output (size / records): 1.4 GiB / 7,844,427
Shuffle read (size / records): 240.1 GiB / 1,961,106,750

Summary metrics for 19250 completed tasks

Metric	Min	25th percentile	Median	75th percentile	Max
Duration	2 s	3 s	3 s	4 s	6 s
GC time	0.5 s	0.7 s	1 s	1 s	2 s
Output (size / records)	45.8 KiB / 374	48.3 KiB / 393	50.3 KiB / 407	52.6 KiB / 421	290.0 KiB / 441
Result serialization time					
Shuffle read (size / records)	11.8 MiB / 93,500	12.4 MiB / 98,250	12.8 MiB / 101,750	13.2 MiB / 105,250	13.8 MiB / 110,250
Shuffle remote reads	11.2 MiB	11.8 MiB	12.1 MiB	12.5 MiB	13.1 MiB
Task deserialization time	5 ms	6 ms	7 ms	13 ms	47 ms

Aggregated metrics by executor (18)

Filter: Filter executors ... 18 executors (all loaded)

Executor ID	Address	Task time	Total tasks	Failed tasks	Succeeded tasks	Output	Shuffle read
1	ip-10-0-0-46.ec2.internal:44057 stderr | stdout	35 s	1199	4	0		
2	ip-10-0-0-198.ec2.internal:45545	10 s	829	4	0		

Lowest cost

Save 75–90% using Reserved Instances and Spot

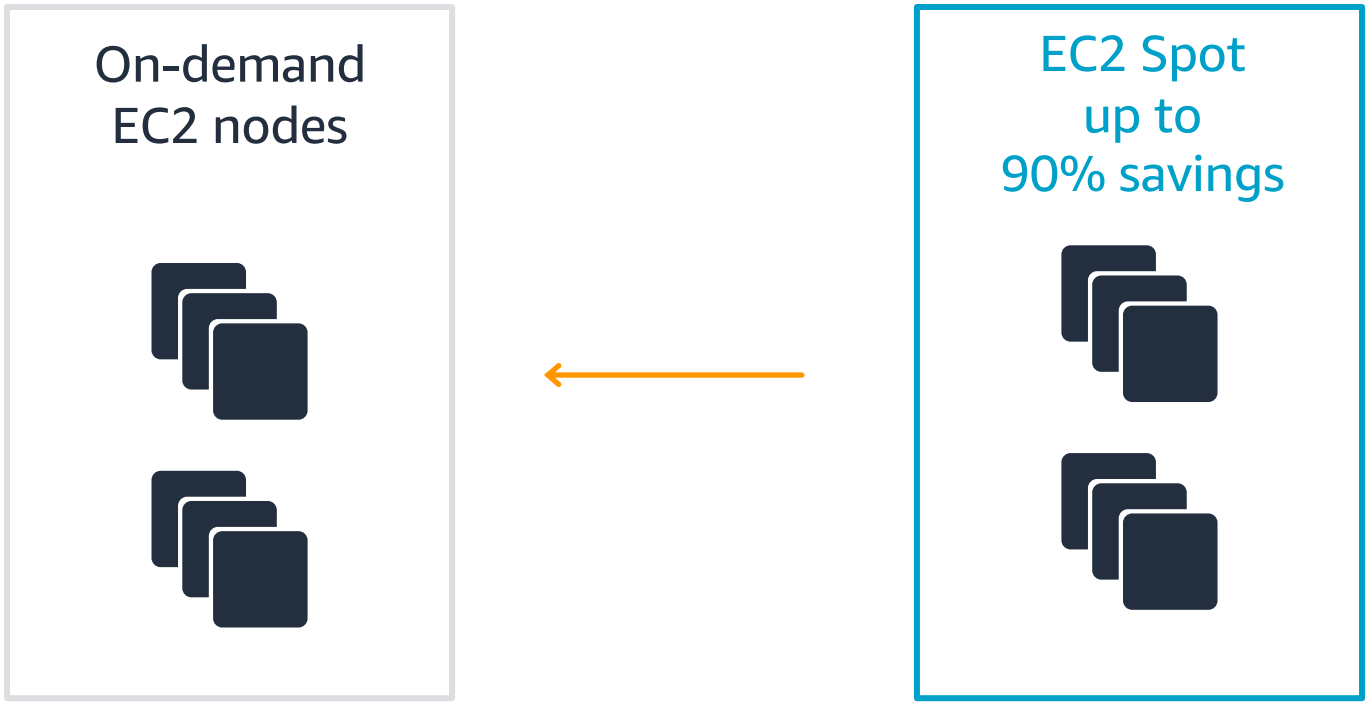

- Commit to a set term and save up to 75% with Reserved Instance
- Run on spare compute capacity and save up to 90% with Spot
  - Pay a fraction of on-demand price
  - Name your bid price. If it exceeds the market, you get the resource
  - Provision from a list of instance types with Spot and on-demand
  - Launch in the most optimal AZ based on capacity/price
  - Spot Block support

aws

Lowest cost

Lowest TCO

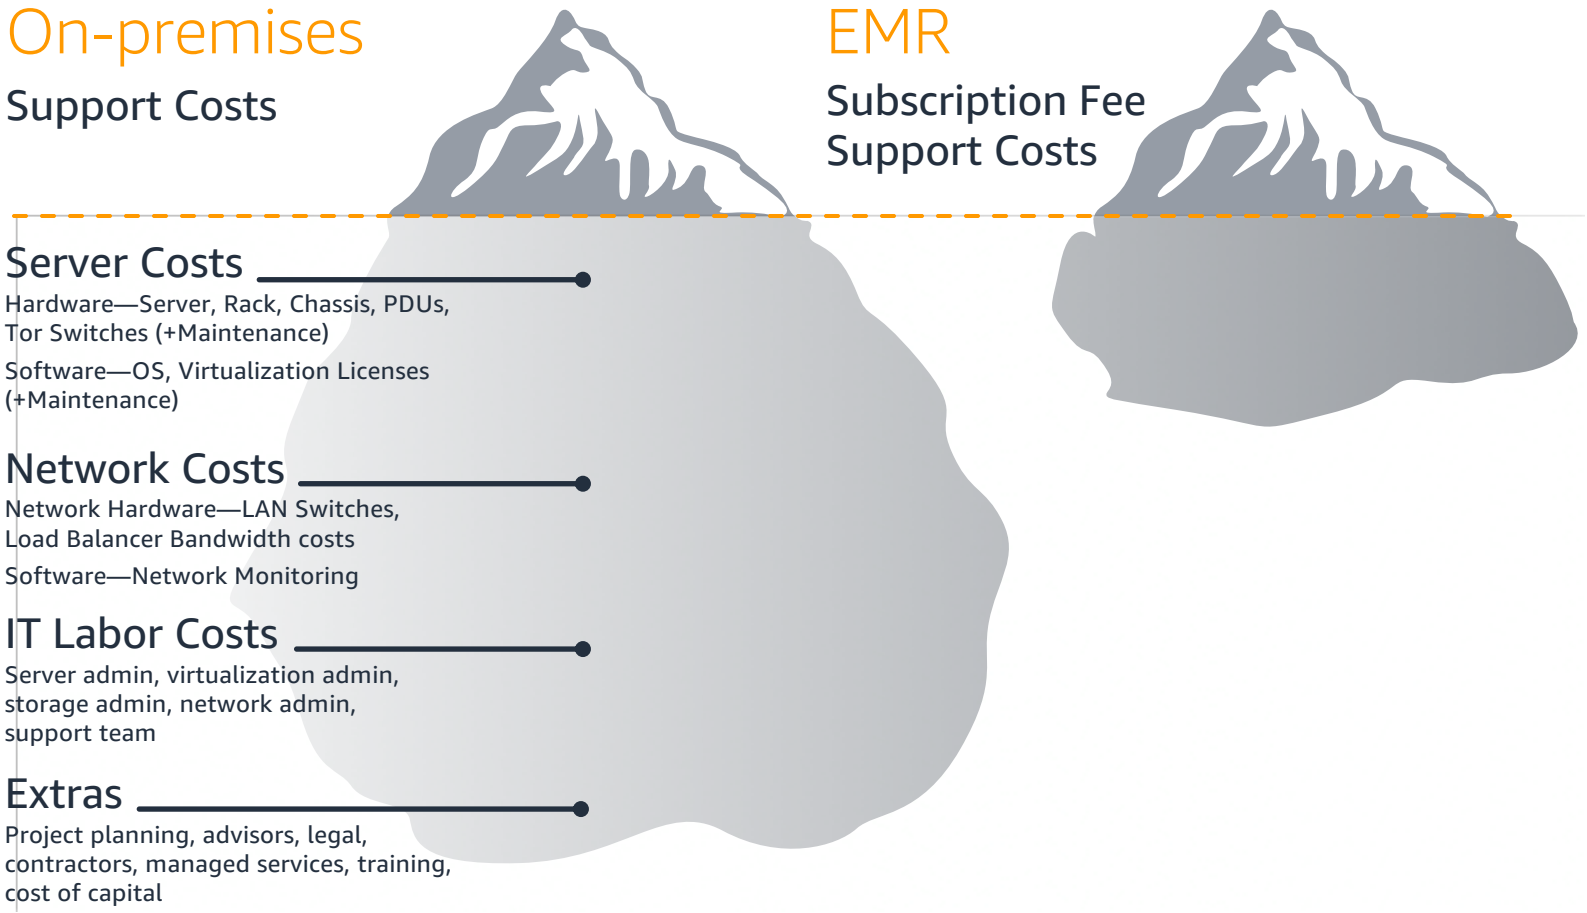

- Less admin time to manage, and support Hadoop clusters
- No up-front costs—hardware acquisition, installation
- Save on operating costs—data center space, power, cooling
- Business Value: Cost of delays, Risk Premium, Competitive abilities, governance, etc.

EC2 Spot and instance fleets

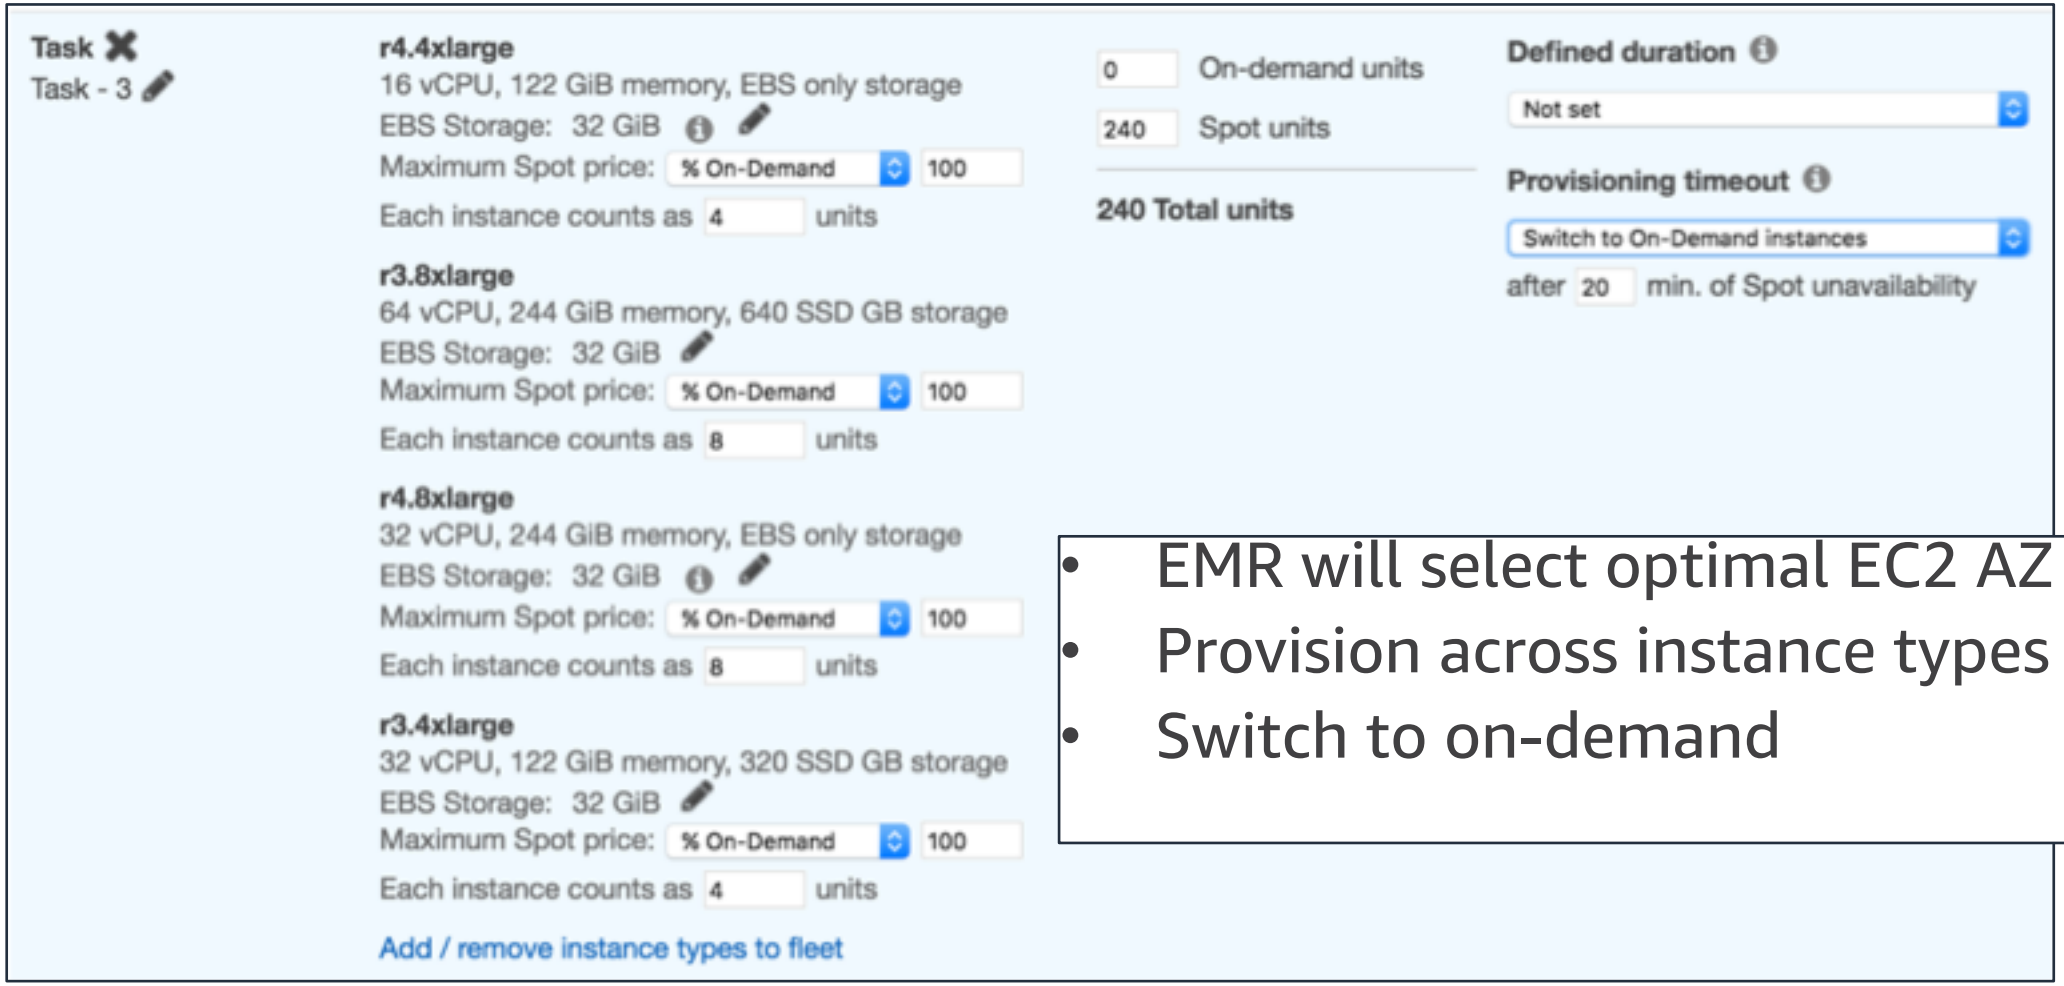

- EMR will select optimal EC2 AZ
- Provision across instance types
- Switch to on-demand

aws

Auto Scaling

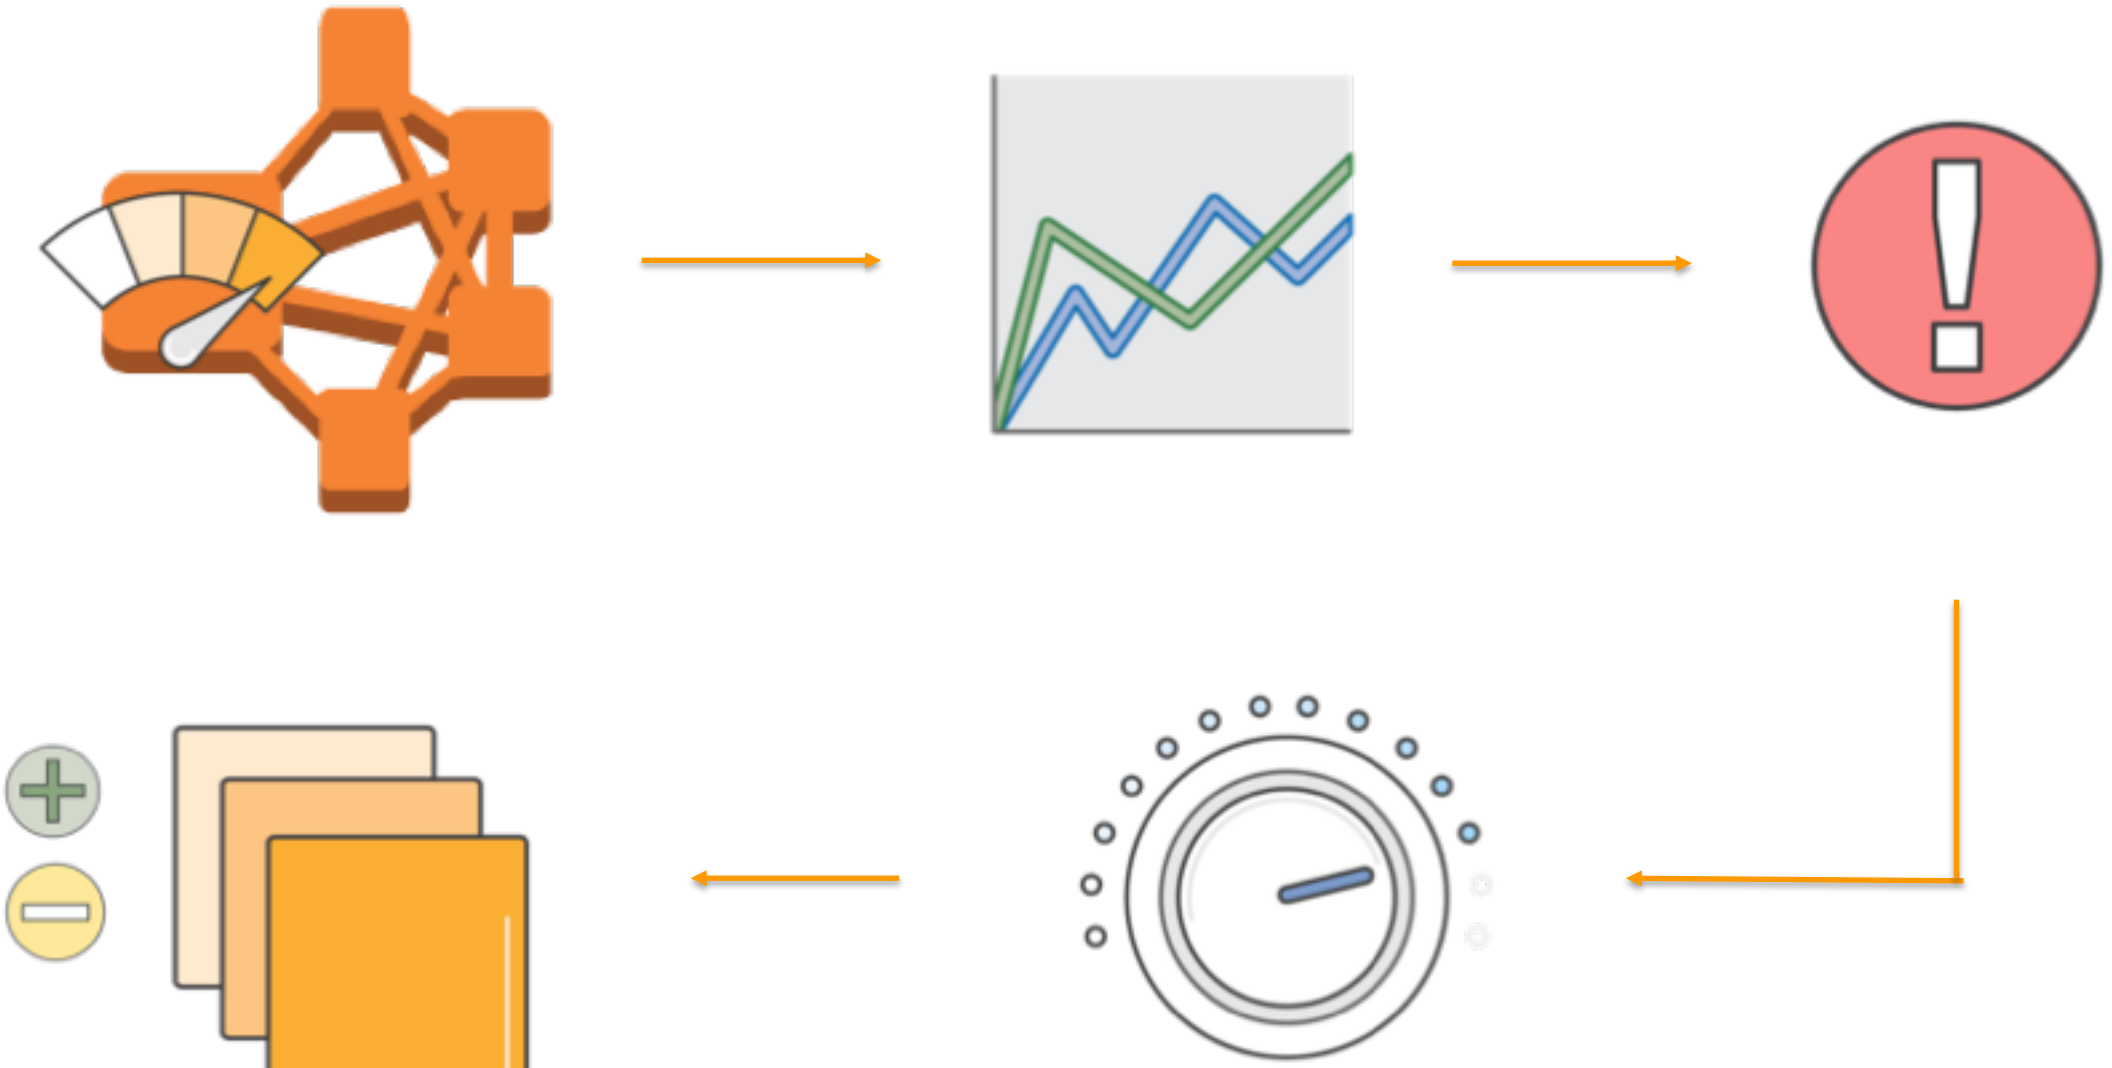

Lowest cost

Hadoop & Spark support included

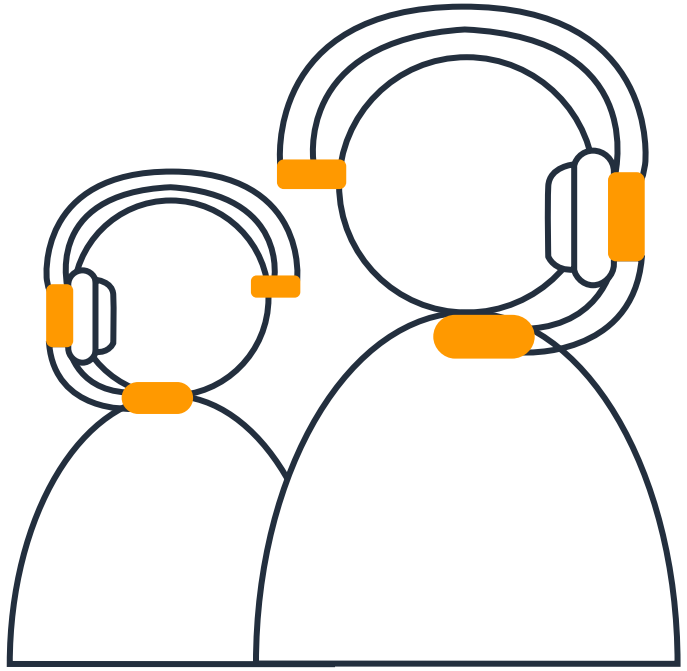

- Hadoop & Spark support included in AWS support
- Support for a 10-node cluster is $4.5K/year
- Up to 1/8th the cost of commercial Hadoop offerings

Customer Story – Zillow Group

Recommendations

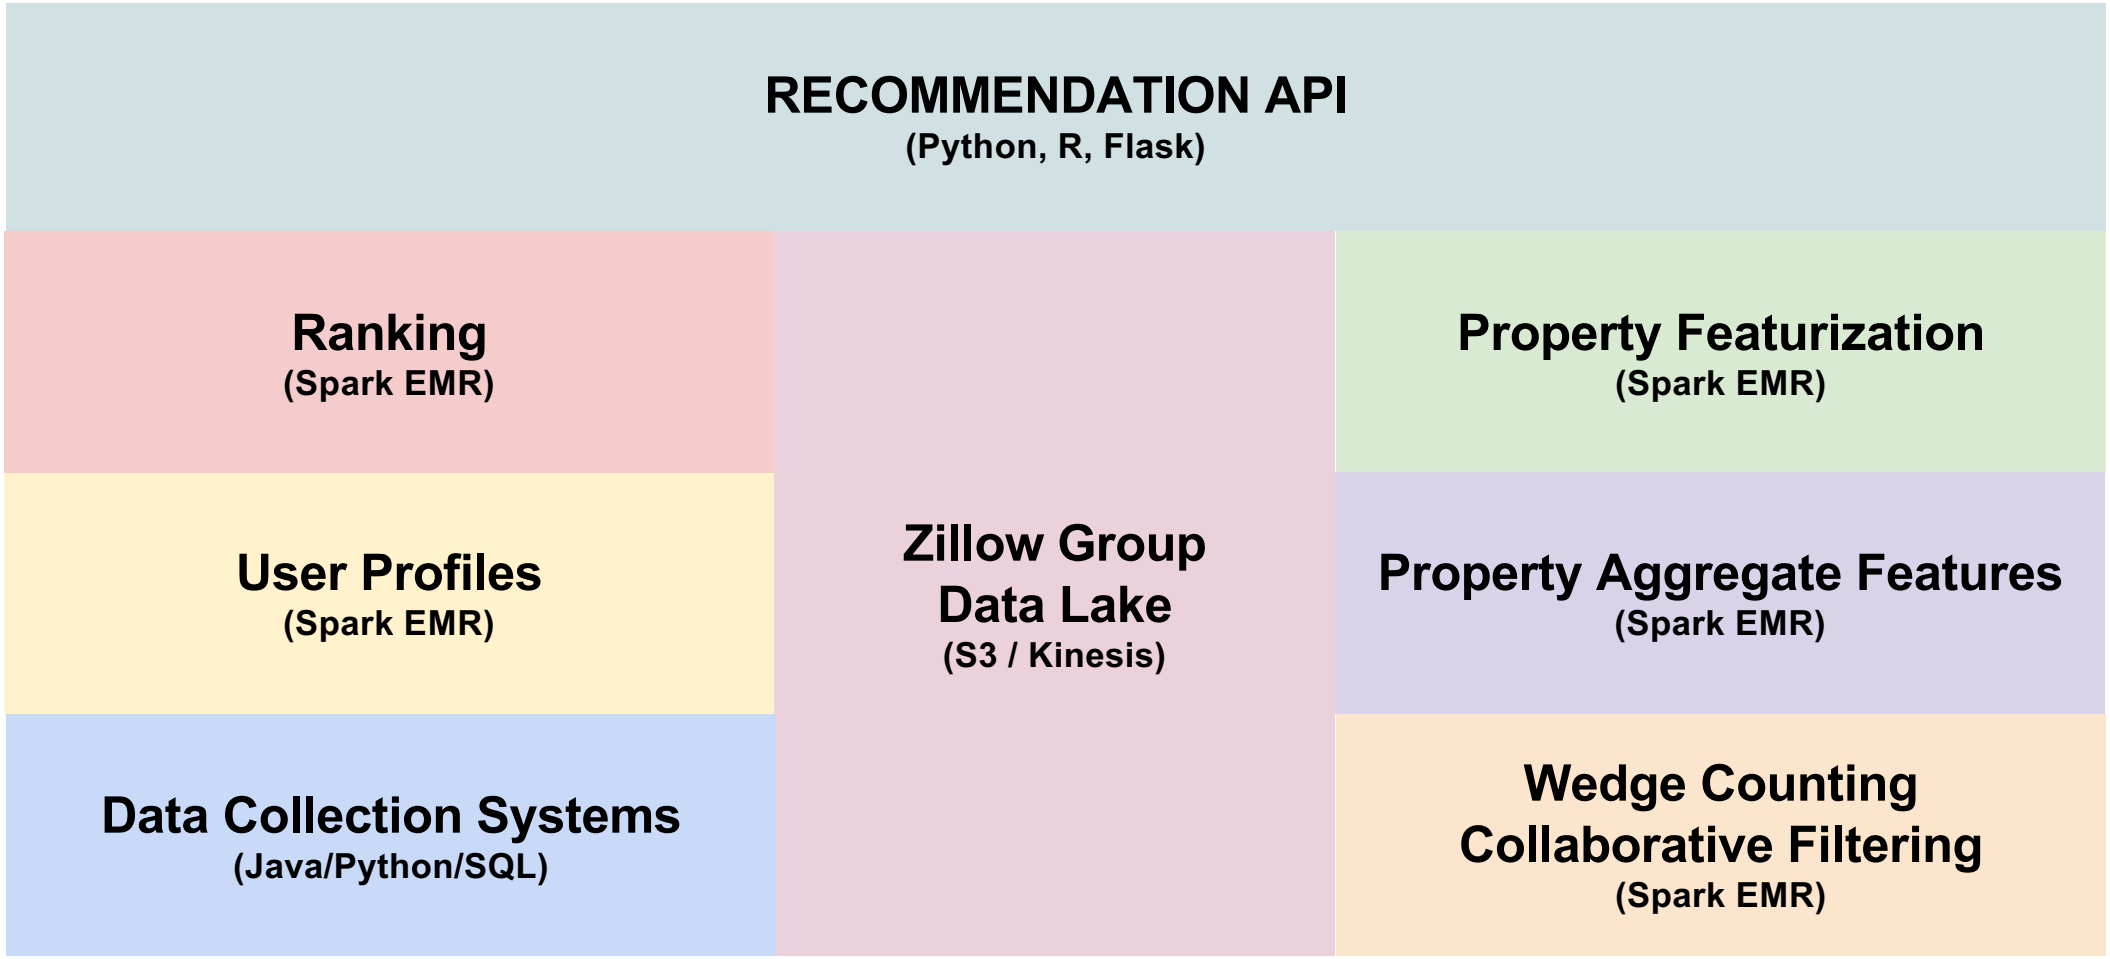

Training & scoring

Zillow GROUP

Collect user behavior and real-estate data, train the various models, generate the candidate set, and make predictions.

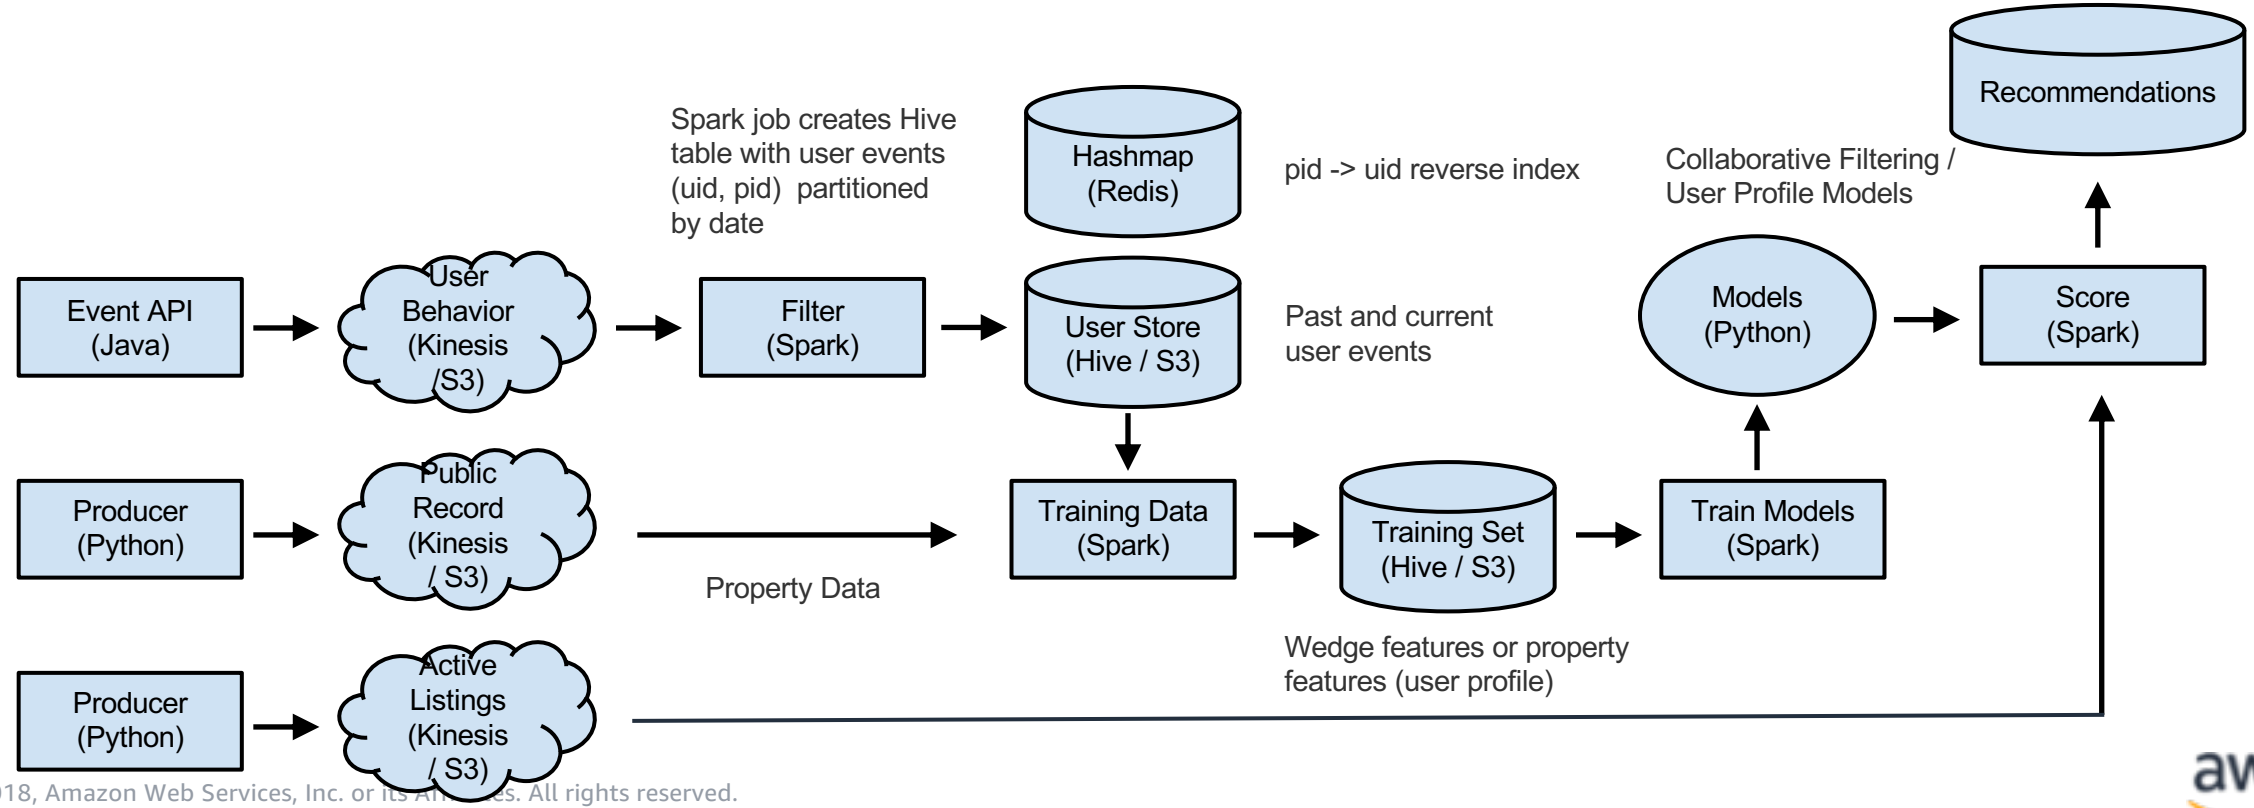

aws

Zestimate

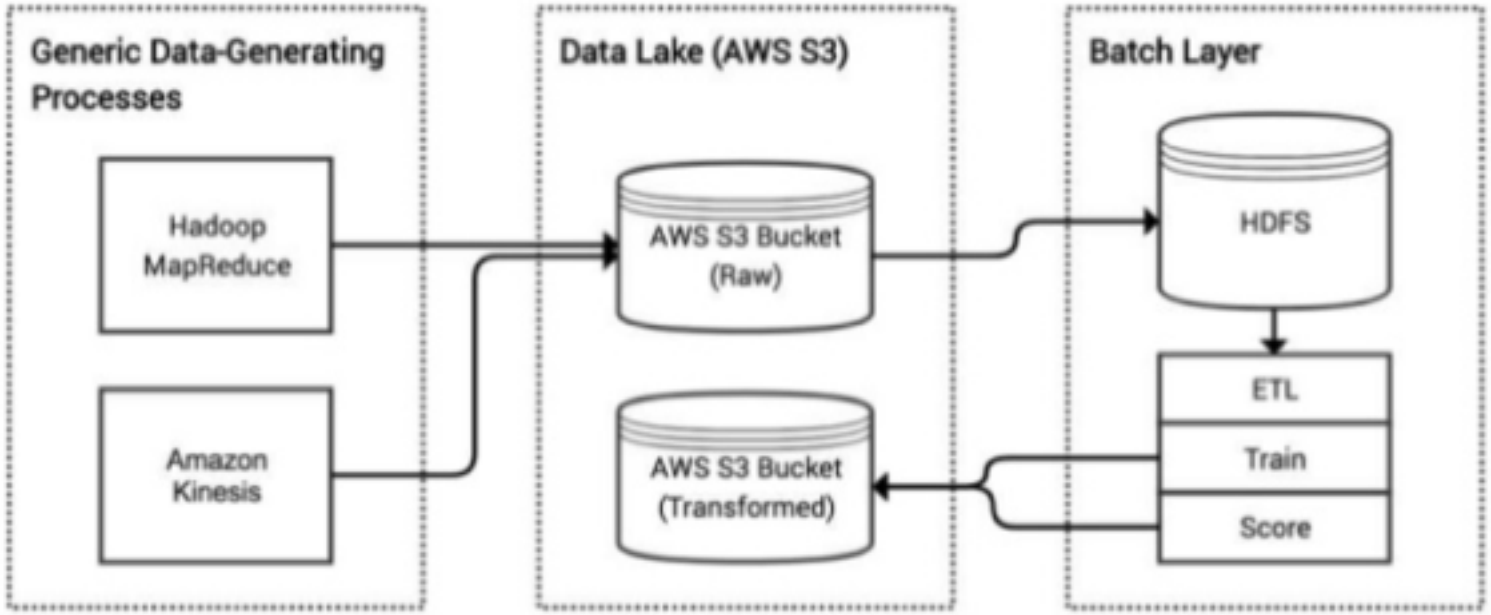

"Previously, given our scale limits using existing proprietary technology, it could take us an entire day or longer to compute a Zestimate," says [Jasjeet] Thind [VP of Data Science and Engineering]. "Now, **we can do it in hours nationwide using Spark on Amazon EMR**, which enables us to quickly perform calculations in parallel on multiple machines."

Ad hoc environment

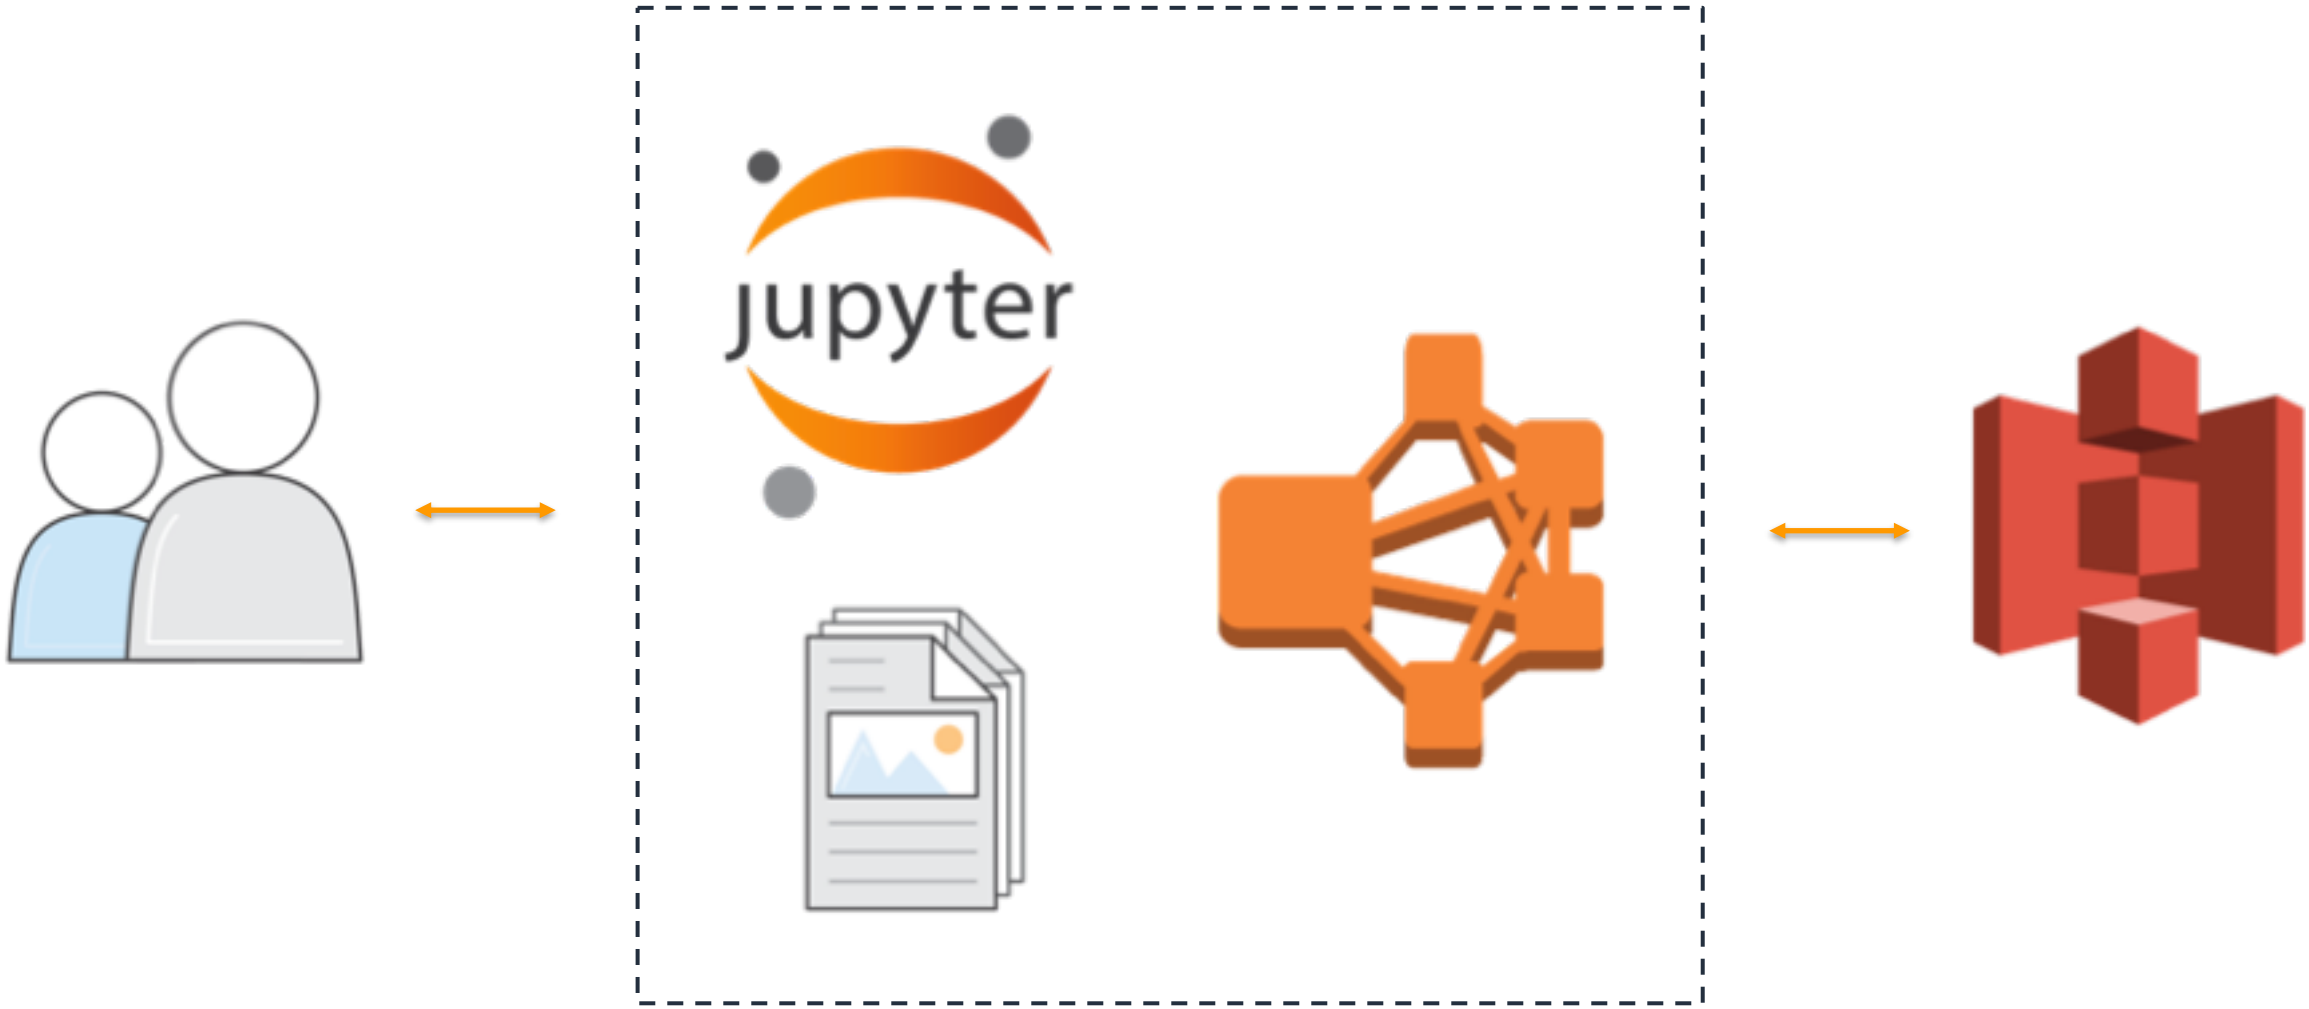

Scale cluster to accommodate more users

Customer Story - DataXu

dataxu

DataXu spun out of MIT Labs to form a PB-scale digital marketing platform.

Challenge:

Deal with billions of impressions, PBs of data. Needed to help world's most valuable brands understand and engage with their consumers (doing real-time advertisement bidding).

Solution:

- Use Spark on EMR and Kinesis to stream and process real-time data from their bidding and retargeting platform
- Land all data in S3 data lake to serve up reports and visualizations

DataXu uses AWS for real-time analytics

dataxu

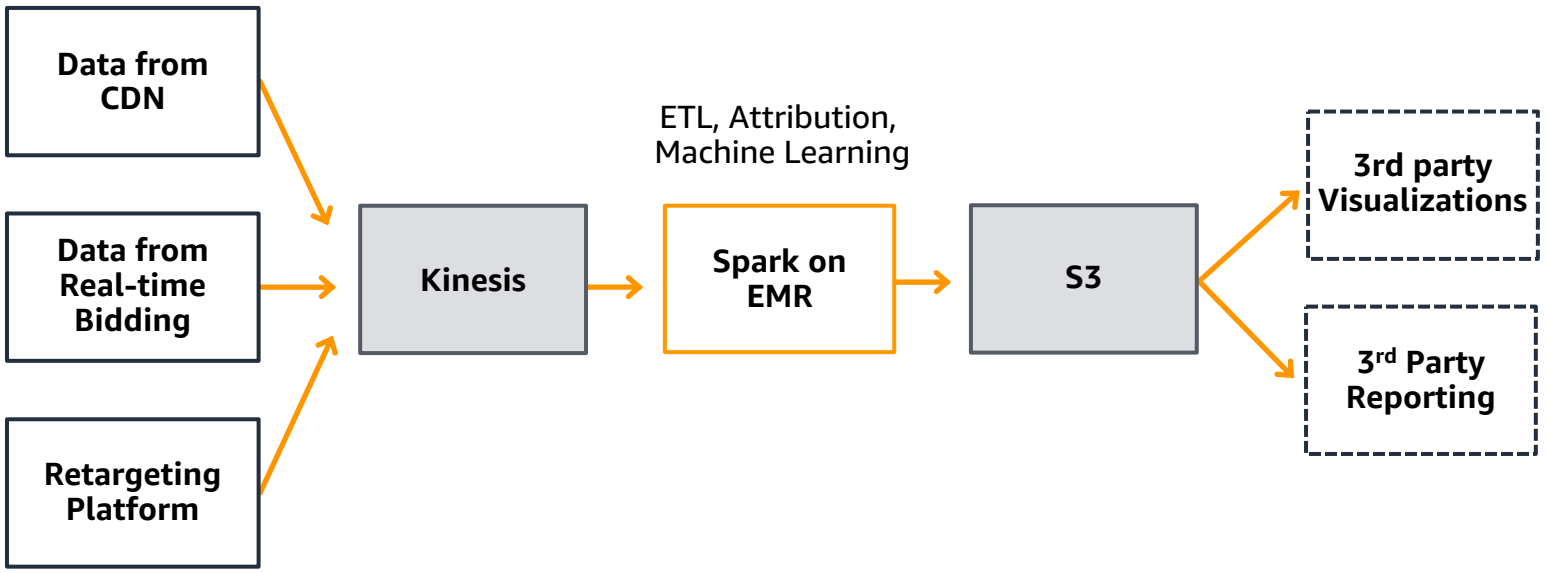

- Stream real-time data from CDN, their real-time bidding platform, and retargeting platform
- Spark on EMR does ETL, attribution, and Machine Learning
- Land data in S3 data lake
- Use third party data visualizations and reporting

aws

DataXu Flows

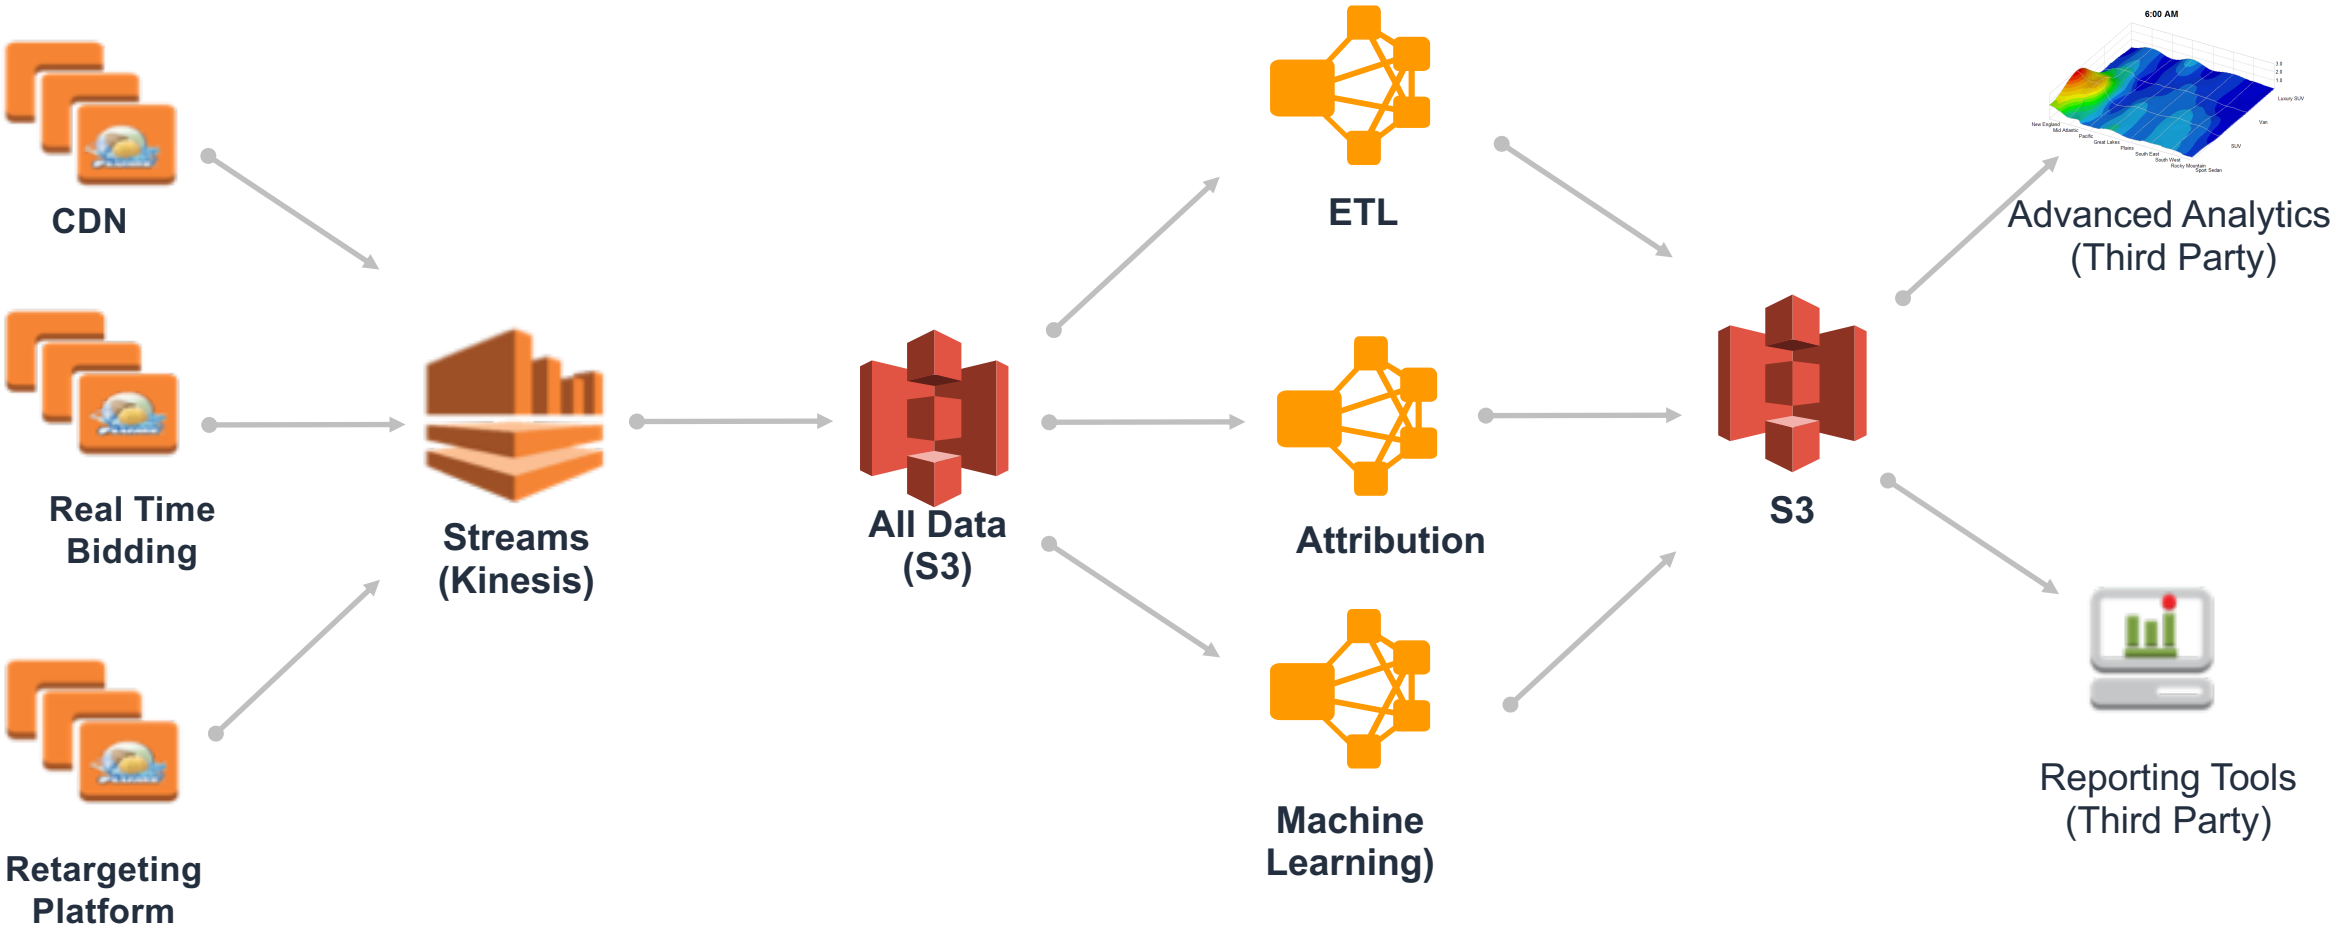

Ecosystem of tools and services

DataXu Spark Pipeline

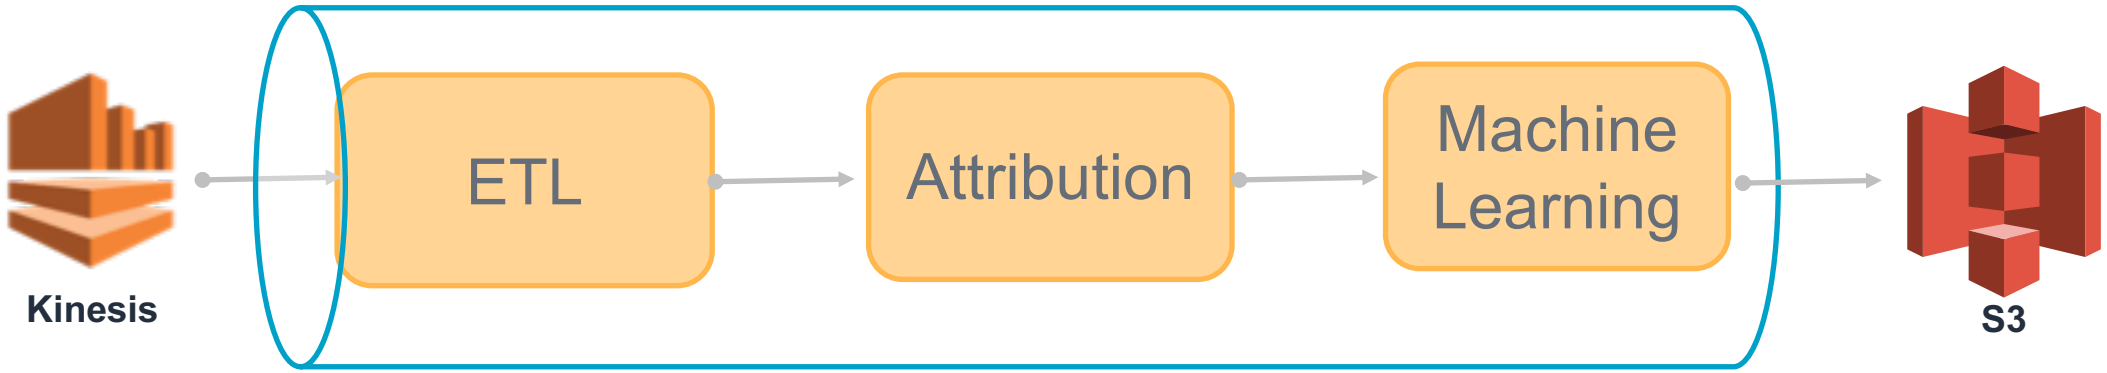

Streaming | Spark SQL | Spark MLib

Parquet, DataSets, DataFrames
Common Backend
Amazon EMR

DataXu Flows – New Generation

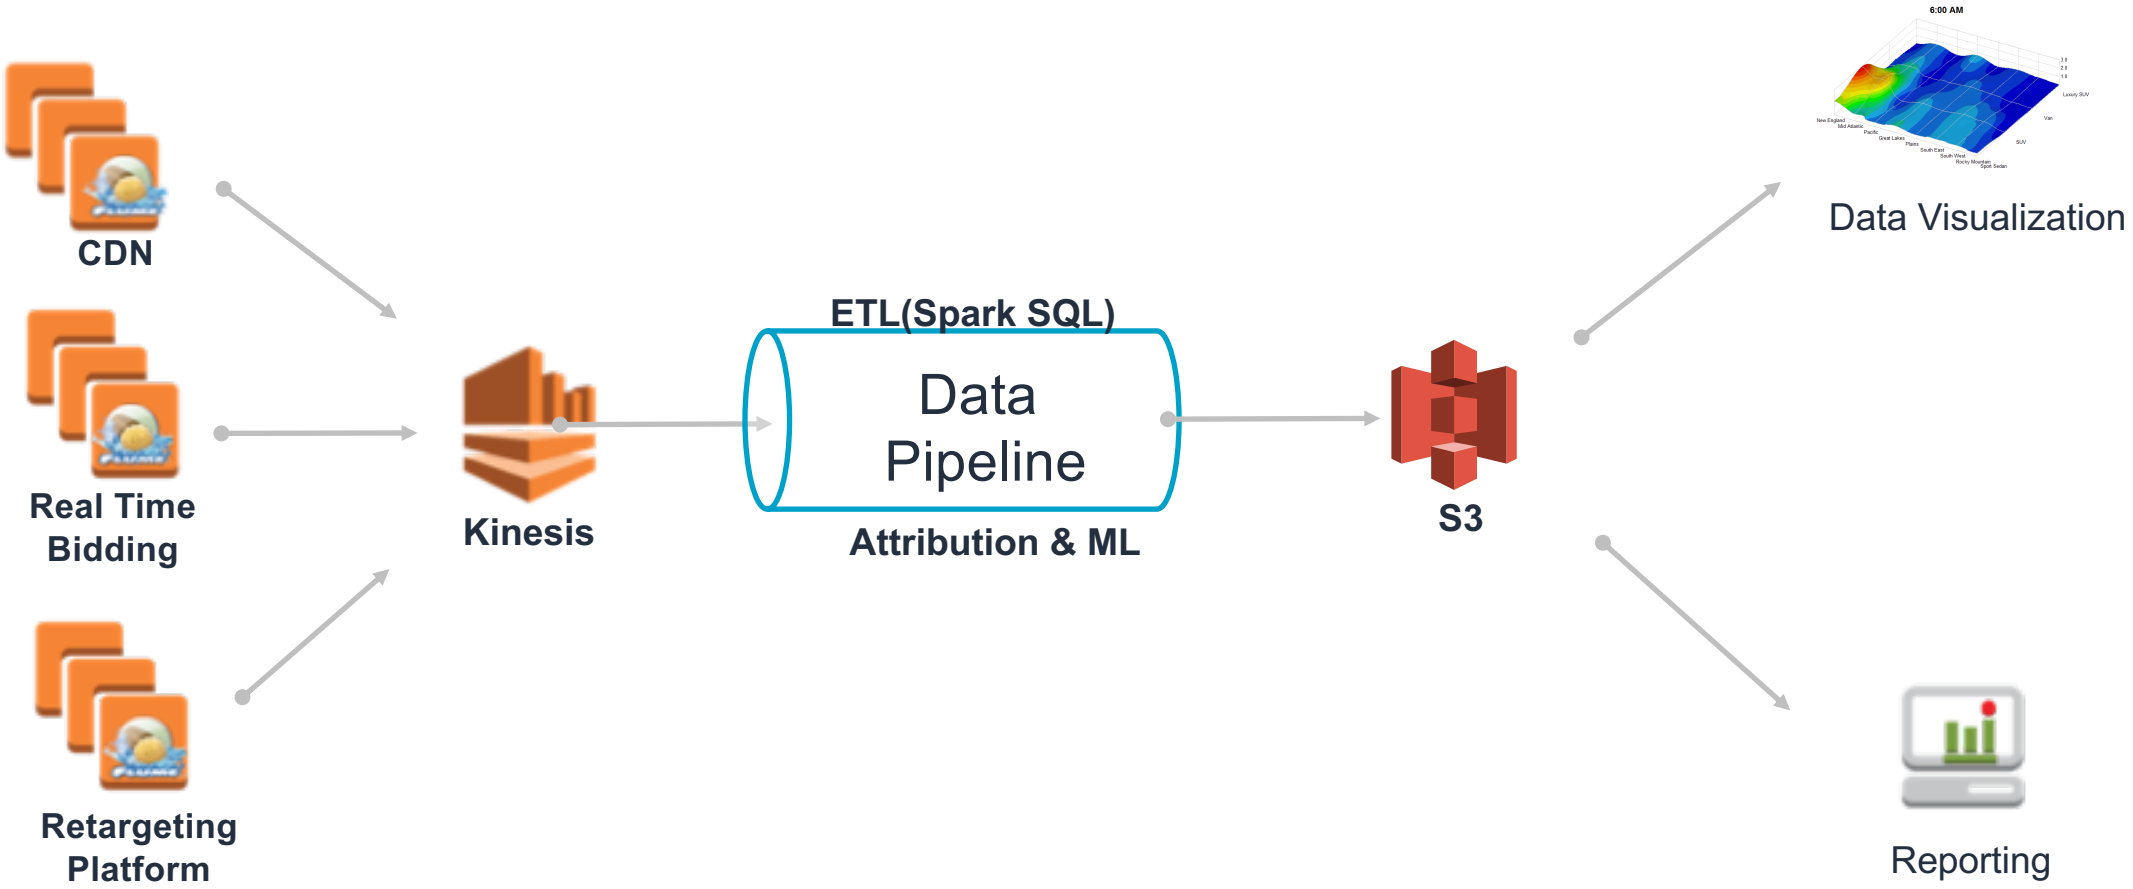

Ecosystem of tools and services

EMR powers the most cloud Hadoop & Spark projects

aws

Thank you!